

04 Amazon Fine Food Reviews Analysis_NaiveBayes

May 29, 2019

1 Amazon Fine Food Reviews Analysis

Data Source: <https://www.kaggle.com/snap/amazon-fine-food-reviews>

EDA: <https://nycdatascience.com/blog/student-works/amazon-fine-foods-visualization/>

The Amazon Fine Food Reviews dataset consists of reviews of fine foods from Amazon.

Number of reviews: 568,454 Number of users: 256,059 Number of products: 74,258 Timespan: Oct 1999 - Oct 2012 Number of Attributes/Columns in data: 10

Attribute Information:

1. Id
2. ProductId - unique identifier for the product
3. UserId - unique identifier for the user
4. ProfileName
5. HelpfulnessNumerator - number of users who found the review helpful
6. HelpfulnessDenominator - number of users who indicated whether they found the review helpful or not
7. Score - rating between 1 and 5
8. Time - timestamp for the review
9. Summary - brief summary of the review
10. Text - text of the review

Objective: Given a review, determine whether the review is positive (rating of 4 or 5) or negative (rating of 1 or 2).

[Q] How to determine if a review is positive or negative? [Ans] We could use Score/Rating. A rating of 4 or 5 can be considered as a positive review. A rating of 1 or 2 can be considered as negative one. A review of rating 3 is considered neutral and such reviews are ignored from our analysis. This is an approximate and proxy way of determining the polarity (positivity/negativity) of a review.

2 [1]. Reading Data

2.1 [1.1] Loading the data

The dataset is available in two forms 1. .csv file 2. SQLite Database

In order to load the data, We have used the SQLITE dataset as it is easier to query the data and visualise the data efficiently.

Here as we only want to get the global sentiment of the recommendations (positive or negative), we will purposefully ignore all Scores equal to 3. If the score is above 3, then the recommendation will be set to "positive". Otherwise, it will be set to "negative".

```
In [1]: %matplotlib inline
import warnings
warnings.filterwarnings("ignore")

import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer

import re
# Tutorial about Python regular expressions: https://pymotw.com/2/re/
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle

from tqdm import tqdm
import os
from joblib import dump, load
from sklearn_pandas import DataFrameMapper
from sklearn.metrics import f1_score, recall_score, precision_score

from sklearn.naive_bayes import MultinomialNB
# importing Cross validation libs
from sklearn.model_selection import train_test_split
from sklearn.model_selection import cross_val_score
from sklearn import model_selection
```

```

# Python script for confusion matrix creation.
from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score
from sklearn.metrics import classification_report
# ROC , AUC curve
# roc curve and auc
from sklearn.datasets import make_classification
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score
from matplotlib import pyplot
from sklearn.metrics import roc_curve, auc
# kFold
from sklearn.model_selection import KFold

from sklearn.model_selection import GridSearchCV

import seaborn as sns

from sklearn.model_selection import TimeSeriesSplit
from sklearn.model_selection import RandomizedSearchCV

In [2]: # using SQLite Table to read data.
con = sqlite3.connect(r'/home/pranay/ML datasource/amazon-fine-food-reviews/database.s

# filtering only positive and negative reviews i.e.
# not taking into consideration those reviews with Score=3
# SELECT * FROM Reviews WHERE Score != 3 LIMIT 500000, will give top 500000 data points.
# you can change the number to any other number based on your computing power

# filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 LIMIT 5
# for tsne assignment you can take 5k data points

filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 LIMIT 100

# Give reviews with Score>3 a positive rating(1), and reviews with a score<3 a negative
def partition(x):
    if x < 3:
        return 0
    return 1

#changing reviews with score less than 3 to be positive and vice-versa
actualScore = filtered_data['Score']
positiveNegative = actualScore.map(partition)
filtered_data['Score'] = positiveNegative
print("Number of data points in our data", filtered_data.shape)
filtered_data.head(3)

Number of data points in our data (100000, 10)

```

```

Out[2]:      Id  ProductId      UserId      ProfileName \
0   1  B001E4KFG0  A3SGXH7AUHU8GW      delmartian
1   2  B00813GRG4  A1D87F6ZCVE5NK      dll pa
2   3  B000LQOCHO  ABXLMWJIXXAIN  Natalia Corres "Natalia Corres"

      HelpfulnessNumerator  HelpfulnessDenominator  Score      Time \
0                        1                        1      1  1303862400
1                        0                        0      0  1346976000
2                        1                        1      1  1219017600

      Summary      Text
0  Good Quality Dog Food  I have bought several of the Vitality canned d...
1    Not as Advertised  Product arrived labeled as Jumbo Salted Peanut...
2  "Delight" says it all  This is a confection that has been around a fe...

```

```

In [3]: display = pd.read_sql_query("""
SELECT UserId, ProductId, ProfileName, Time, Score, Text, COUNT(*)
FROM Reviews
GROUP BY UserId
HAVING COUNT(*)>1
""", con)

```

```

In [4]: print(display.shape)
display.head()

```

```

(80668, 7)

```

```

Out[4]:      UserId      ProductId      ProfileName      Time      Score \
0  #oc-R115TNMSPFT9I7  B005ZBZLT4      Breyton  1331510400      2
1  #oc-R11D9D7SHXIJB9  B005HG9ESG  Louis E. Emory "hoppy"  1342396800      5
2  #oc-R11DNU2NBKQ23Z  B005ZBZLT4      Kim Cieszykowski  1348531200      1
3  #oc-R1105J5ZVQE25C  B005HG9ESG      Penguin Chick  1346889600      5
4  #oc-R12KPBODL2B5ZD  B007OSBEV0  Christopher P. Presta  1348617600      1

      Text      COUNT(*)
0  Overall its just OK when considering the price...      2
1  My wife has recurring extreme muscle spasms, u...      3
2  This coffee is horrible and unfortunately not ...      2
3  This will be the bottle that you grab from the...      3
4  I didnt like this coffee. Instead of telling y...      2

```

```

In [5]: display[display['UserId']=='AZY10LLTJ71NX']

```

```

Out[5]:      UserId      ProductId      ProfileName      Time \
80638  AZY10LLTJ71NX  B001ATMQK2  undertheshrine "undertheshrine"  1296691200

      Score      Text      COUNT(*)
80638      5  I bought this 6 pack because for the price tha...      5

```

```
In [6]: display['COUNT(*)'].sum()
```

```
Out[6]: 393063
```

3 [2] Exploratory Data Analysis

3.1 [2.1] Data Cleaning: Deduplication

It is observed (as shown in the table below) that the reviews data had many duplicate entries. Hence it was necessary to remove duplicates in order to get unbiased results for the analysis of the data. Following is an example:

```
In [7]: display= pd.read_sql_query("""
        SELECT *
        FROM Reviews
        WHERE Score != 3 AND UserId="AR5J8UI46CURR"
        ORDER BY ProductID
        """, con)
display.head()
```

```
Out[7]:
```

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator \
0	78445	B000HDL1RQ	AR5J8UI46CURR	Geetha Krishnan	2
1	138317	B000HDOPYC	AR5J8UI46CURR	Geetha Krishnan	2
2	138277	B000HDOPYM	AR5J8UI46CURR	Geetha Krishnan	2
3	73791	B000HDOPZG	AR5J8UI46CURR	Geetha Krishnan	2
4	155049	B000PAQ75C	AR5J8UI46CURR	Geetha Krishnan	2

	HelpfulnessDenominator	Score	Time \
0	2	5	1199577600
1	2	5	1199577600
2	2	5	1199577600
3	2	5	1199577600
4	2	5	1199577600

	Summary \
0	LOACKER QUADRATINI VANILLA WAFERS
1	LOACKER QUADRATINI VANILLA WAFERS
2	LOACKER QUADRATINI VANILLA WAFERS
3	LOACKER QUADRATINI VANILLA WAFERS
4	LOACKER QUADRATINI VANILLA WAFERS

	Text
0	DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
1	DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
2	DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
3	DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
4	DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...

As it can be seen above that same user has multiple reviews with same values for HelpfulnessNumerator, HelpfulnessDenominator, Score, Time, Summary and Text and on doing analysis it was found that ProductId=B000HDOPZG was Loacker Quadratini Vanilla Wafer Cookies, 8.82-Ounce Packages (Pack of 8) ProductId=B000HDL1RQ was Loacker Quadratini Lemon Wafer Cookies, 8.82-Ounce Packages (Pack of 8) and so on

It was inferred after analysis that reviews with same parameters other than ProductId belonged to the same product just having different flavour or quantity. Hence in order to reduce redundancy it was decided to eliminate the rows having same parameters.

The method used for the same was that we first sort the data according to ProductId and then just keep the first similar product review and delete the others. for eg. in the above just the review for ProductId=B000HDL1RQ remains. This method ensures that there is only one representative for each product and deduplication without sorting would lead to possibility of different representatives still existing for the same product.

```
In [8]: #Sorting data according to ProductId in ascending order
sorted_data=filtered_data.sort_values('ProductId', axis=0, ascending=True, inplace=False)
```

```
In [9]: #Deduplication of entries
final=sorted_data.drop_duplicates(subset={"UserId","ProfileName","Time","Text"}, keep='first')
final.shape
```

```
Out[9]: (87775, 10)
```

```
In [10]: #Checking to see how much % of data still remains
(final['Id'].size*1.0)/(filtered_data['Id'].size*1.0)*100
```

```
Out[10]: 87.775
```

Observation:- It was also seen that in two rows given below the value of HelpfulnessNumerator is greater than HelpfulnessDenominator which is not practically possible hence these two rows too are removed from calculations

```
In [11]: display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND Id=44737 OR Id=64422
ORDER BY ProductID
""", con)

display.head()
```

```
Out[11]:
```

	Id	ProductId	UserId	ProfileName	\
0	64422	B000MIDROQ	A161DK06JJMCYF	J. E. Stephens	"Jeanne"
1	44737	B001EQ55RW	A2V0I904FH7ABY		Ram

	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	\
0	3	1	5	1224892800	
1	3	2	4	1212883200	

```

                                Summary \
0          Bought This for My Son at College
1 Pure cocoa taste with crunchy almonds inside

                                Text
0 My son loves spaghetti so I didn't hesitate or...
1 It was almost a 'love at first bite' - the per...

```

```
In [12]: final=final[final.HelpfulnessNumerator<=final.HelpfulnessDenominator]
```

```
In [13]: #Before starting the next phase of preprocessing lets see the number of entries left
        print(final.shape)
```

```

#How many positive and negative reviews are present in our dataset?
final['Score'].value_counts()

```

```
(87773, 10)
```

```

Out[13]: 1    73592
         0    14181
         Name: Score, dtype: int64

```

4 [3] Preprocessing

4.1 [3.1]. Preprocessing Review Text

Now that we have finished deduplication our data requires some preprocessing before we go on further with analysis and making the prediction model.

Hence in the Preprocessing phase we do the following in the order below:-

1. Begin by removing the html tags
2. Remove any punctuations or limited set of special characters like , or . or # etc.
3. Check if the word is made up of english letters and is not alpha-numeric
4. Check to see if the length of the word is greater than 2 (as it was researched that there is no adjective in 2-letters)
5. Convert the word to lowercase
6. Remove Stopwords
7. Finally Snowball Stemming the word (it was observed to be better than Porter Stemming)

After which we collect the words used to describe positive and negative reviews

```

In [14]: # printing some random reviews
        sent_0 = final['Text'].values[0]
        print(sent_0)
        print("="*50)

        sent_1000 = final['Text'].values[1000]
        print(sent_1000)

```

```

print("="*50)

sent_1500 = final['Text'].values[1500]
print(sent_1500)
print("="*50)

sent_4900 = final['Text'].values[4900]
print(sent_4900)
print("="*50)

```

```

My dogs loves this chicken but its a product from China, so we wont be buying it anymore.  Its
=====
The Candy Blocks were a nice visual for the Lego Birthday party but the candy has little taste
=====
was way to hot for my blood, took a bite and did a jig  lol
=====
My dog LOVES these treats. They tend to have a very strong fish oil smell. So if you are afraid
=====

```

```

In [15]: # remove urls from text python: https://stackoverflow.com/a/40823105/4084039
sent_0 = re.sub(r"http\S+", "", sent_0)
sent_1000 = re.sub(r"http\S+", "", sent_1000)
sent_1500 = re.sub(r"http\S+", "", sent_1500)
sent_4900 = re.sub(r"http\S+", "", sent_4900)

print(sent_0)

```

```

My dogs loves this chicken but its a product from China, so we wont be buying it anymore.  Its

```

```

In [16]: # https://stackoverflow.com/questions/16206380/python-beautifulsoup-how-to-remove-all
from bs4 import BeautifulSoup

soup = BeautifulSoup(sent_0, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_1000, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_1500, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

```



```
soup = BeautifulSoup(sent_4900, 'lxml')
text = soup.get_text()
print(text)
```

My dogs loves this chicken but its a product from China, so we wont be buying it anymore. Its
=====

The Candy Blocks were a nice visual for the Lego Birthday party but the candy has little taste
=====

was way to hot for my blood, took a bite and did a jig lol
=====

My dog LOVES these treats. They tend to have a very strong fish oil smell. So if you are afraid

```
In [17]: # https://stackoverflow.com/a/47091490/4084039
import re
```

```
def decontracted(phrase):
    # specific
    phrase = re.sub(r"won't", "will not", phrase)
    phrase = re.sub(r"can't", "can not", phrase)

    # general
    phrase = re.sub(r"n't", " not", phrase)
    phrase = re.sub(r"\ 're", " are", phrase)
    phrase = re.sub(r"\ 's", " is", phrase)
    phrase = re.sub(r"\ 'd", " would", phrase)
    phrase = re.sub(r"\ 'll", " will", phrase)
    phrase = re.sub(r"\ 't", " not", phrase)
    phrase = re.sub(r"\ 've", " have", phrase)
    phrase = re.sub(r"\ 'm", " am", phrase)
    return phrase
```

```
In [18]: sent_1500 = decontracted(sent_1500)
print(sent_1500)
print("="*50)
```

was way to hot for my blood, took a bite and did a jig lol
=====

```
In [19]: #remove words with numbers python: https://stackoverflow.com/a/18082370/4084039
sent_0 = re.sub(r"\S*\d\S*", "", sent_0).strip()
print(sent_0)
```

My dogs loves this chicken but its a product from China, so we wont be buying it anymore. Its

```
In [20]: #remove spacial character: https://stackoverflow.com/a/5843547/4084039
sent_1500 = re.sub(r'[^A-Za-z0-9]+', ' ', sent_1500)
print(sent_1500)
```

was way to hot for my blood took a bite and did a jig lol

```
In [21]: # https://gist.github.com/sebleier/554280
# we are removing the words from the stop words list: 'no', 'nor', 'not'
# <br /><br /> ==> after the above steps, we are getting "br br"
# we are including them into stop words list
# instead of <br /> if we have <br/> these tags would have reumoved in the 1st step

stopwords= set(['br', 'the', 'i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselv
"you'll", "you'd", 'your', 'yours', 'yourself', 'yourselves', 'he', 'him'
'she', "she's", 'her', 'hers', 'herself', 'it', "it's", 'its', 'itself',
'theirs', 'themselves', 'what', 'which', 'who', 'whom', 'this', 'that', "t
'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'have', 'has', 'l
'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or', 'because', 'as
'at', 'by', 'for', 'with', 'about', 'against', 'between', 'into', 'through
'above', 'below', 'to', 'from', 'up', 'down', 'in', 'out', 'on', 'off', 'o
'then', 'once', 'here', 'there', 'when', 'where', 'why', 'how', 'all', 'an
'most', 'other', 'some', 'such', 'only', 'own', 'same', 'so', 'than', 'too
's', 't', 'can', 'will', 'just', 'don', "don't", 'should', "should've", 'r
've', 'y', 'ain', 'aren', "aren't", 'couldn', "couldn't", 'didn', "didn't
"hadn't", 'hasn', "hasn't", 'haven', "haven't", 'isn', "isn't", 'ma', 'mi
"mustn't", 'needn', "needn't", 'shan', "shan't", 'shouldn', "shouldn't",
'won', "won't", 'wouldn', "wouldn't"])
```

```
In [22]: #filtered out whole reviews
```

```
from bs4 import BeautifulSoup
# Combining all the above stundents
from tqdm import tqdm
# tqdm is for printing the status bar
word_counter = []
def filterised_text(text):
    preprocessed_text = []
    for sentence in tqdm(text):
        sentence = re.sub(r"http\S+", "", sentence)
        sentence = BeautifulSoup(sentence, 'lxml').get_text()
        sentence = decontracted(sentence)
        sentence = re.sub("\S*\d\S*", "", sentence).strip()
        sentence = re.sub('[^A-Za-z]+', ' ', sentence)
        # https://gist.github.com/sebleier/554280
        sentence = ' '.join(e.lower() for e in sentence.split() if e.lower() not in s
        count = len(sentence.split())
        word_counter.append(count)
        preprocessed_text.append(sentence.strip())
    return preprocessed_text
```

```
In [23]: preprocessed_reviews = filterised_text(final['Text'].values)
```

```
final['preprocessed_reviews'] = preprocessed_reviews
preprocessed_reviews[1822]
```

100%|| 87773/87773 [00:29<00:00, 2956.18it/s]

Out[23]: 'taste great using air popper not great little seeds fall popping'

```
In [24]: final['numbers_of_words'] = word_counter
word_counter[1822]
```

Out[24]: 11

4.2 Preprocessing Review Summary

```
In [25]: preprocessed_summary = filterised_text(final['Summary'].values)
final['preprocessed_summary'] = preprocessed_summary
preprocessed_summary[1822]
```

100%|| 87773/87773 [00:18<00:00, 4810.09it/s]

Out[25]: 'pop corn'

4.2.1 Splitting data

We have considered 100 k points

```
In [26]: X = final['preprocessed_reviews']
y = final['Score']

# split the data set into train and test
X_train, x_test, y_train, y_test = model_selection.train_test_split(X, y, test_size=0

print(X_train.shape, x_test.shape, y_train.shape, y_test.shape)
```

(61441,) (26332,) (61441,) (26332,)

5 [4] Featurization

5.1 [4.1] BAG OF WORDS

```
In [27]: ##BoW
```

```
count_vect = CountVectorizer(ngram_range=(1,2), min_df=10) #in scikit-learn

# train data
X_train_bow = count_vect.fit_transform(X_train)
```

```

# test data
x_test_bow = count_vect.transform(x_test)

print('X_train_bow',X_train_bow.shape)
print('==='*10)
print('x_test_bow',x_test_bow.shape)

X_train_bow (61441, 36487)
=====
x_test_bow (26332, 36487)

```

6 [5] Assignment 4: Apply Naive Bayes

```

In [28]: from sklearn.naive_bayes import MultinomialNB
# importing Cross validation libs
from sklearn.model_selection import train_test_split
from sklearn.model_selection import cross_val_score
from sklearn import model_selection

# Python script for confusion matrix creation.
from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score
from sklearn.metrics import classification_report
# ROC , AUC curve
# roc curve and auc
from sklearn.datasets import make_classification
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score
from matplotlib import pyplot
from sklearn.metrics import roc_curve, auc
# kFold
from sklearn.model_selection import KFold

from sklearn.model_selection import GridSearchCV

import seaborn as sns

from sklearn.model_selection import TimeSeriesSplit
from sklearn.model_selection import RandomizedSearchCV

In [29]: # Common Methods

# define the range, through which we are going to find alpha-hyperparameter
# alpha_values = (1e-4, 1e-3,1e-2,1e-1, 1e0,1e1)

```

```

# alpha_values = np.geomspace(1e-3, 10)

alpha_values = (1e-4, 1e-3, 1e-2, 0.05, 1e-1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1e0, 3, 5, 7, 9)

def finding_best_alpha(X_tr, y_tr):
    # instantiate a Multinomial Naive Bayes model
    nb = MultinomialNB()

    param_grid = dict(alpha=alpha_values)
    print(param_grid)

    # For time based splitting
    tscv = TimeSeriesSplit(n_splits=15)

    # instantiate the training grid search model
    train_grid = GridSearchCV(nb, param_grid, cv=tscv, scoring='roc_auc', n_jobs=-1, verbose=0)

    # fit the training data to train model
    train_grid.fit(X_tr, y_tr)

    return train_grid

# plot a graph which show difference between validation error and training error

def plotAccuracyGraph(training_grid):
    alpha_range = [i for i in alpha_values]
    accuracy = [i for i in training_grid.cv_results_['mean_train_score']]
    accuracy_test = [i for i in training_grid.cv_results_['mean_test_score']]
    plt.semilogx(alpha_range, accuracy, 'r', label='train_accuracy')
    plt.semilogx(alpha_range, accuracy_test, 'b', label='validation_accuracy')
    plt.title('Accuracy plot')
    plt.xlabel('Alpha')
    plt.ylabel('Accuracy')
    plt.grid('on')
    plt.legend()
    plt.show()

# https://www.geeksforgeeks.org/confusion-matrix-machine-learning/
def plotConfusionMatrix(y_test, pred):
    # calculate confusion matrix
    cm = confusion_matrix(y_test, pred)
    class_label = ['negative', 'positive']
    df_conf_matrix = pd.DataFrame(cm, index=class_label, columns=class_label)
    # heatmap --> Plot rectangular data as a color-encoded matrix.
    sns.heatmap(df_conf_matrix, annot=True, fmt='d')
    # give title to graph
    plt.title("Confusion Matrix")
    # mention axis label

```

```

plt.xlabel("Predicted")
plt.ylabel("Actual")
# show the plot
plt.show()

# https://machinelearningmastery.com/roc-curves-and-precision-recall-curves-for-class
# plot AUC curve
def plotAUC_ROC(nb_optimal,X_train, y_train, X_test, y_test):
    # predict probabilities
    test_probs = nb_optimal.predict_proba(X_test)
    train_probs = nb_optimal.predict_proba(X_train)

    # keep probabilities for the positive outcome only
    test_probs = test_probs[:, 1]
    train_probs = train_probs[:, 1]

    # calculate AUC
    test_auc = roc_auc_score(y_test, test_probs)
    train_auc = roc_auc_score(y_train, train_probs)

    # calculate roc curve
    train_fpr, train_tpr, thresholds = roc_curve(y_train, train_probs)
    test_fpr, test_tpr, thresholds2 = roc_curve(y_test, test_probs)

    # plot no skill
    pyplot.plot([0, 1], [0, 1], linestyle='--')

    # plot the roc curve for the model
    pyplot.plot(train_fpr, train_tpr, 'r',marker='.', label="train AUC =" +str(train_auc))
    pyplot.plot(test_fpr, test_tpr, 'b',marker='.',label="test AUC =" +str(test_auc))

    pyplot.legend()
    pyplot.xlabel("K: hyperparameter")
    pyplot.ylabel("AUC")
    pyplot.title("ERROR PLOTS")
    # show the plot
    pyplot.show()

    return train_auc, test_auc

class color:
    PURPLE = '\033[95m'
    CYAN = '\033[96m'
    DARKCYAN = '\033[36m'
    BLUE = '\033[94m'
    GREEN = '\033[92m'
    YELLOW = '\033[93m'

```

```

RED = '\033[91m'
BOLD = '\033[1m'
UNDERLINE = '\033[4m'
END = '\033[0m'

```

```

# https://stackoverflow.com/questions/11116697/how-to-get-most-informative-features-f
def important_features(feature_names, classifier, n=20):
    class_labels = classifier.classes_
    topn_class1 = sorted(zip(classifier.feature_count_[0], feature_names), reverse=True)
    topn_class2 = sorted(zip(classifier.feature_count_[1], feature_names), reverse=True)
    print(color.BOLD+"Important words in negative reviews"+color.END)
    print('\n'+color.BOLD+'\t Class Label '+color.END, class_labels[0])
    for coef, feat in topn_class1:
        print('{:.3f}'.format(coef), '\t'+feat)
    print("-----\n")
    print(color.BOLD+"Important words in positive reviews"+color.END)
    print('\n'+color.BOLD+'\t Class Label '+color.END, class_labels[1])
    for coef, feat in topn_class2:
        print('{:.3f}'.format(coef), '\t'+feat)

```

7 Applying Multinomial Naive Bayes

7.1 [5.1] Applying Naive Bayes on BOW

In [30]: `bow_train = finding_best_alpha(X_train_bow, y_train)`

```

# view the complete results (list of named tuples)
print("=====Training=====")
print (bow_train.best_score_)
print (bow_train.best_params_)
print (bow_train.best_estimator_)

```

```

plotAccuracyGraph(bow_train)

```

{'alpha': (0.0001, 0.001, 0.01, 0.05, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1.0, 3, 5, ...
Fitting 15 folds for each of 21 candidates, totalling 315 fits

```

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.
[Parallel(n_jobs=-1)]: Done   5 tasks      | elapsed:    1.5s
[Parallel(n_jobs=-1)]: Done  10 tasks      | elapsed:    1.7s
[Parallel(n_jobs=-1)]: Done  17 tasks      | elapsed:    1.9s
[Parallel(n_jobs=-1)]: Done  24 tasks      | elapsed:    2.0s
[Parallel(n_jobs=-1)]: Done  33 tasks      | elapsed:    2.3s
[Parallel(n_jobs=-1)]: Done  42 tasks      | elapsed:    2.6s
[Parallel(n_jobs=-1)]: Done  53 tasks      | elapsed:    2.8s

```

```
[Parallel(n_jobs=-1)]: Batch computation too fast (0.1969s.) Setting batch_size=2.
[Parallel(n_jobs=-1)]: Done 64 tasks      | elapsed: 3.1s
[Parallel(n_jobs=-1)]: Done 86 tasks      | elapsed: 3.8s
[Parallel(n_jobs=-1)]: Done 112 tasks     | elapsed: 4.5s
[Parallel(n_jobs=-1)]: Done 142 tasks     | elapsed: 5.2s
[Parallel(n_jobs=-1)]: Done 172 tasks     | elapsed: 5.9s
[Parallel(n_jobs=-1)]: Done 206 tasks     | elapsed: 6.8s
[Parallel(n_jobs=-1)]: Done 240 tasks     | elapsed: 7.6s
[Parallel(n_jobs=-1)]: Done 278 tasks     | elapsed: 8.5s
[Parallel(n_jobs=-1)]: Done 308 out of 315 | elapsed: 9.3s remaining: 0.2s
[Parallel(n_jobs=-1)]: Done 315 out of 315 | elapsed: 9.5s finished
```


=====Training=====

0.9336869731761391

{'alpha': 0.5}

MultinomialNB(alpha=0.5, class_prior=None, fit_prior=True)

/home/pranay/anaconda3/lib/python3.7/site-packages/matplotlib/cbook/__init__.py:424: MatplotlibDeprecationWarning: Passing one of 'on', 'true', 'off', 'false' as a boolean is deprecated; use an actual boolean value instead.
warn_deprecated("2.2", "Passing one of 'on', 'true', 'off', 'false' as a "




```

In [31]: optimal_alpha = bow_train.best_params_.get('alpha')

nb_optimal= MultinomialNB(alpha=optimal_alpha, class_prior=None, fit_prior=True)

# fitting the model
nb_optimal.fit(X_train_bow,y_train)

# predict the response
test_pred = nb_optimal.predict(x_test_bow)
train_pred = nb_optimal.predict(X_train_bow)

print('\n'+color.RED+'Optimal best alpha is : '+color.END+color.BOLD+str(optimal_alpha))

# plot confusion matrix

print('\n'+color.BOLD +'Confusion Matrix for Train data'+color.END)
plotConfusionMatrix(y_train,train_pred)

print('\n'+color.BOLD +'Confusion Matrix for Test data'+color.END)
plotConfusionMatrix(y_test,test_pred)

# plot AUC
train_auc,test_auc = plotAUC_ROC(nb_optimal,X_train_bow, y_train,x_test_bow, y_test )
print('\n'+color.RED+'AUC (Train): '+color.END+color.BOLD+str(train_auc)+color.END)
print('\n'+color.RED+'AUC (Test): '+color.END+color.BOLD+str(test_auc)+color.END)

# f1 score
score = f1_score(y_test,test_pred)
print('\n'+color.RED+'F1 SCORE (Train) : '+color.END+color.BOLD+str(f1_score(y_train,train_pred))+color.END)
print('\n'+color.RED+'F1 SCORE (Test) : '+color.END+color.BOLD+str(score)+color.END)


# recall
recall = metrics.recall_score(y_test, test_pred)
print('\n'+color.RED+'RECALL (Train): '+color.END+color.BOLD+str(metrics.recall_score(y_train,train_pred))+color.END)
print('\n'+color.RED+'RECALL (Test): '+color.END+color.BOLD+str(recall)+color.END)

# precision
precision = metrics.precision_score(y_test, test_pred)
print('\n'+color.RED+'PRECISION (Train) : '+color.END+color.BOLD+str(metrics.precision_score(y_train,train_pred))+color.END)
print('\n'+color.RED+'PRECISION (Test) : '+color.END+color.BOLD+str(precision)+color.END)


```


Optimal best alpha is : 0.5

Confusion Matrix for Train data



Confusion Matrix for Test data





AUC (Train): 0.9744685792725499

AUC (Test): 0.9450283466889644

F1 SCORE (Train) : 0.9626646928201332

F1 SCORE (Test) : 0.9495393922727898

RECALL (Train): 0.947230422622747

RECALL (Test): 0.9372964169381107

PRECISION (Train) : 0.9786102694784999

PRECISION (Test) : 0.9621064363332404

7.1.1 [5.1.1] Top 10 important features of positive class

7.1.2 [5.1.2] Top 10 important features of negative class

```
In [32]: important_features(count_vect.get_feature_names(), nb_optimal, 10)
```

Important words in negative reviews

	Class Label 0
16486.000	not
5374.000	like
4246.000	would
4098.000	product
4052.000	taste
3322.000	one
2651.000	coffee
2627.000	good
2498.000	flavor
2489.000	no

Important words in positive reviews

	Class Label 1
48321.000	not
21545.000	like
18904.000	good
17323.000	great
15165.000	one
13878.000	taste
13426.000	coffee
12529.000	flavor
12454.000	love
12439.000	would

7.1.3 Feature Engineering

Till now we only consider Text review as feature, we are adding some extra feature like **review summary** and **number of words** in review and test our model improves efficiency or not.

We have considered on 50000 points due to memory issue.

```
In [33]: # https://sondosatwi.wordpress.com/2017/08/01/using-text-data-and-dataframemapper-in-
```

```
X = final[:50000]
y = final['Score'][:50000]
```

```
# split the data set into train and test
```

```
X_train, x_test, y_train, y_test = model_selection.train_test_split(X, y, test_size=0
```

```

print(X_train.shape, x_test.shape, y_train.shape, y_test.shape)

mapper = DataFrameMapper([
    ('preprocessed_reviews', CountVectorizer(ngram_range=(1,3), min_df=10)),
    ('preprocessed_summary', CountVectorizer(ngram_range=(1,3), min_df=10)),
    ('numbers_of_words', None),
])

train_features = mapper.fit_transform(X_train)
test_features = mapper.transform(x_test)

optimal_alpha = bow_train.best_params_.get('alpha')

nb_optimal= MultinomialNB(alpha=optimal_alpha, class_prior=None, fit_prior=True)

# fitting the model
nb_optimal.fit(train_features,y_train)

# predict the response
test_pred = nb_optimal.predict(test_features)
train_pred = nb_optimal.predict(train_features)

# plot confusion matrix

print('\n'+color.BOLD + 'Confusion Matrix for Train data'+color.END)
plotConfusionMatrix(y_train,train_pred)

print('\n'+color.BOLD + 'Confusion Matrix for Test data'+color.END)
plotConfusionMatrix(y_test,test_pred)

# plot AUC
train_auc,test_auc = plotAUC_ROC(nb_optimal,train_features, y_train,test_features, y_test)
print('\n'+color.RED+'AUC (Train): '+color.END+color.BOLD+str(train_auc)+color.END)
print('\n'+color.RED+'AUC (Test): '+color.END+color.BOLD+str(test_auc)+color.END)

# f1 score
score = f1_score(y_test,test_pred)
print('\n'+color.RED+'F1 SCORE (Train) : '+color.END+color.BOLD+str(f1_score(y_train,train_pred))+color.END)
print('\n'+color.RED+'F1 SCORE (Test) : '+color.END+color.BOLD+str(score)+color.END)

# recall
recall = metrics.recall_score(y_test, test_pred)
print('\n'+color.RED+'RECALL (Train): '+color.END+color.BOLD+str(metrics.recall_score(y_train,train_pred))+color.END)
print('\n'+color.RED+'RECALL (Test): '+color.END+color.BOLD+str(recall)+color.END)


# precision
precision = metrics.precision_score(y_test, test_pred)

```



```
print('\n'+color.RED+'PRECISION (Train) : '+color.END+color.BOLD+str(metrics.precision_train))
print('\n'+color.RED+'PRECISION (Test) : '+color.END+color.BOLD+str(precision)+color.END)
```

(35000, 13) (15000, 13) (35000,) (15000,)

Confusion Matrix for Train data



Confusion Matrix for Test data



AUC (Train): 0.9816129910282342

AUC (Test): 0.9566592879401167

F1 SCORE (Train) : 0.9684253972621587

F1 SCORE (Test) : 0.9546202303093528

RECALL (Train): 0.9527612344342177

RECALL (Test): 0.9424813817144668

PRECISION (Train) : 0.9846132326199468

PRECISION (Test) : 0.9670758474920739

7.1.4 [5.1.1] Top 10 important features of positive class

7.1.5 [5.1.2] Top 10 important features of negative class

```
In [34]: merged_features_vectorizer = mapper.features[0][1].get_feature_names() + mapper.features[1][1].get_feature_names()
         important_features(merged_features_vectorizer, nb_optimal, 10)
```

Important words in negative reviews

	Class Label 0
8654.000	not
2765.000	like
2208.000	product
2190.000	would
2009.000	taste
1808.000	one
1421.000	good
1379.000	no
1379.000	food
1330.000	not

Important words in positive reviews

	Class Label 1
26431.000	not
11388.000	like
10595.000	good
9743.000	great
8923.000	tea
8144.000	one

7372.000	taste
6762.000	love
6690.000	flavor
6674.000	product

As we can see that by consider 'review text', 'summary text' and 'number of words' in review text all together, **AUC value is changed from 0.94502 to 0.95665**

7.2 [4.2] TF-IDF

```
In [35]: X = final['preprocessed_reviews']
        y = final['Score']

        # split the data set into train and test
        X_train, x_test, y_train, y_test = model_selection.train_test_split(X, y, test_size=0

        print(X_train.shape, x_test.shape, y_train.shape, y_test.shape)

        tf_idf_vect = TfidfVectorizer(ngram_range=(1,3), min_df=10) #in scikit-learn

        # train data
        X_train_tfidf = tf_idf_vect.fit_transform(X_train)

        # test data
        x_test_tfidf = tf_idf_vect.transform(x_test)

        print('X_train_tfidf', X_train_tfidf.shape)
        print('=='*10)
        print('x_test_tfidf', x_test_tfidf.shape)

(61441,) (26332,) (61441,) (26332,)
X_train_tfidf (61441, 40217)
=====
x_test_tfidf (26332, 40217)
```

7.3 [5.2] Applying Naive Bayes on TFIDF

```
In [36]: tfidf_train = finding_best_alpha(X_train_tfidf,y_train)

        # view the complete results (list of named tuples)
        print("====Training=====")
        print (tfidf_train.best_score_)
        print (tfidf_train.best_params_)
        print (tfidf_train.best_estimator_)

        plotAccuracyGraph(tfidf_train)
```

```
{'alpha': (0.0001, 0.001, 0.01, 0.05, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1.0, 3, 5, ...  
Fitting 15 folds for each of 21 candidates, totalling 315 fits
```


```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.  
[Parallel(n_jobs=-1)]: Batch computation too fast (0.0550s.) Setting batch_size=6.  
[Parallel(n_jobs=-1)]: Done   5 tasks      | elapsed:    0.1s  
[Parallel(n_jobs=-1)]: Done  20 tasks      | elapsed:    0.7s  
[Parallel(n_jobs=-1)]: Done  62 tasks      | elapsed:    1.9s  
[Parallel(n_jobs=-1)]: Done 104 tasks      | elapsed:    2.9s  
[Parallel(n_jobs=-1)]: Done 158 tasks      | elapsed:    4.1s  
[Parallel(n_jobs=-1)]: Done 212 tasks      | elapsed:    5.3s  
[Parallel(n_jobs=-1)]: Done 278 tasks      | elapsed:    7.0s  
[Parallel(n_jobs=-1)]: Done 315 out of 315 | elapsed:    7.9s finished  
/home/pranay/anaconda3/lib/python3.7/site-packages/matplotlib/cbook/__init__.py:424: Matplotlib  
Passing one of 'on', 'true', 'off', 'false' as a boolean is deprecated; use an actual boolean  
warn_deprecated("2.2", "Passing one of 'on', 'true', 'off', 'false' as a "
```

```
=====Training=====
```

```
0.945188332274678
```

```
{'alpha': 0.1}
```

```
MultinomialNB(alpha=0.1, class_prior=None, fit_prior=True)
```



```

In [37]: optimal_alpha = tfidf_train.best_params_.get('alpha')

print('\n'+color.RED+'Optimal best alpha is : '+color.END+color.BOLD+str(optimal_alpha))

nb_optimal= MultinomialNB(alpha=optimal_alpha, class_prior=None, fit_prior=True)

# fitting the model
nb_optimal.fit(X_train_tfidf,y_train)

# predict the response
test_pred = nb_optimal.predict(x_test_tfidf)
train_pred = nb_optimal.predict(X_train_tfidf)

# plot confusion matrix

print('\n'+color.BOLD +'Confusion Matrix for Train data'+color.END)
plotConfusionMatrix(y_train,train_pred)

print('\n'+color.BOLD +'Confusion Matrix for Test data'+color.END)
plotConfusionMatrix(y_test,test_pred)

# plot AUC
train_auc,test_auc = plotAUC_ROC(nb_optimal,X_train_tfidf, y_train,x_test_tfidf, y_test)
print('\n'+color.RED+'AUC (Train): '+color.END+color.BOLD+str(train_auc)+color.END)
print('\n'+color.RED+'AUC (Test): '+color.END+color.BOLD+str(test_auc)+color.END)

# f1 score
score = f1_score(y_test,test_pred)
print('\n'+color.RED+'F1 SCORE (Train) : '+color.END+color.BOLD+str(f1_score(y_train,train_pred))+color.END)
print('\n'+color.RED+'F1 SCORE (Test) : '+color.END+color.BOLD+str(score)+color.END)


# recall
recall = metrics.recall_score(y_test, test_pred)
print('\n'+color.RED+'RECALL (Train): '+color.END+color.BOLD+str(metrics.recall_score(y_train,train_pred))+color.END)
print('\n'+color.RED+'RECALL (Test): '+color.END+color.BOLD+str(recall)+color.END)

# precision
precision = metrics.precision_score(y_test, test_pred)
print('\n'+color.RED+'PRECISION (Train) : '+color.END+color.BOLD+str(metrics.precision_score(y_train,train_pred))+color.END)
print('\n'+color.RED+'PRECISION (Test) : '+color.END+color.BOLD+str(precision)+color.END)


```


Optimal best alpha is : 0.1

Confusion Matrix for Train data



Confusion Matrix for Test data





AUC (Train): 0.9846151250529447

AUC (Test): 0.9559963012212672

F1 SCORE (Train) : 0.9661914460285133

F1 SCORE (Test) : 0.9496387220336033

RECALL (Train): 0.9904832193909261

RECALL (Test): 0.9870159247195078

PRECISION (Train) : 0.9430626698966289

PRECISION (Test) : 0.9149890957892971

7.3.1 [5.2.1] Top 10 important features of positive class

7.3.2 [5.2.2] Top 10 important features of negative class

```
In [38]: important_features(tf_idf_vect.get_feature_names(), nb_optimal, 10)
```

Important words in negative reviews

	Class Label 0
519.547	not
228.356	like
212.892	product
203.388	would
202.627	taste
162.527	coffee
150.752	one
132.950	flavor
130.494	no
117.575	good

Important words in positive reviews

	Class Label 1
1456.513	not
1030.260	great
972.931	good
923.692	like
882.647	coffee
817.438	love
808.220	tea
729.114	one
719.204	taste
712.064	product

7.3.3 Feature Engineering

Till now we only consider Text review as feature, we are adding some extra feature like **review summary** and **number of words** in review and test our model improves efficiency or not.

We have considered on 50000 points due to memory issue.

```
In [41]: # https://sondosatwi.wordpress.com/2017/08/01/using-text-data-and-dataframemapper-in-
```

```
X = final[:50000]
y = final['Score'][:50000]
```

```
# split the data set into train and test
```

```
X_train, x_test, y_train, y_test = model_selection.train_test_split(X, y, test_size=0
```

```

print(X_train.shape, x_test.shape, y_train.shape, y_test.shape)

mapper = DataFrameMapper([
    ('preprocessed_reviews', TfidfVectorizer(ngram_range=(1,2), min_df=10)),
    ('preprocessed_summary', TfidfVectorizer(ngram_range=(1,2), min_df=10)),
    ('numbers_of_words', None),
])

train_features = mapper.fit_transform(X_train)
test_features = mapper.transform(x_test)

optimal_alpha = bow_train.best_params_.get('alpha')

optimal_alpha

nb_optimal= MultinomialNB(alpha=optimal_alpha, class_prior=None, fit_prior=True)

# fitting the model
nb_optimal.fit(train_features,y_train)

# predict the response
test_pred = nb_optimal.predict(test_features)
train_pred = nb_optimal.predict(train_features)

# plot confusion matrix

print('\n'+color.BOLD + 'Confusion Matrix for Train data'+color.END)
plotConfusionMatrix(y_train,train_pred)

print('\n'+color.BOLD + 'Confusion Matrix for Test data'+color.END)
plotConfusionMatrix(y_test,test_pred)

# plot AUC
train_auc,test_auc = plotAUC_ROC(nb_optimal,train_features, y_train,test_features, y_
print('\n'+color.RED+'AUC (Train): '+color.END+color.BOLD+str(train_auc)+color.END)
print('\n'+color.RED+'AUC (Test): '+color.END+color.BOLD+str(test_auc)+color.END)

# f1 score
score = f1_score(y_test,test_pred)
print('\n'+color.RED+'F1 SCORE (Train) : '+color.END+color.BOLD+str(f1_score(y_train,
print('\n'+color.RED+'F1 SCORE (Test) : '+color.END+color.BOLD+str(score)+color.END)

# recall
recall = metrics.recall_score(y_test, test_pred)
print('\n'+color.RED+'RECALL (Train): '+color.END+color.BOLD+str(metrics.recall_score
print('\n'+color.RED+'RECALL (Test): '+color.END+color.BOLD+str(recall)+color.END)


```

```



# precision
precision = metrics.precision_score(y_test, test_pred)
print('\n'+color.RED+'PRECISION (Train) : '+color.END+color.BOLD+str(metrics.precision))
print('\n'+color.RED+'PRECISION (Test) : '+color.END+color.BOLD+str(precision)+color.END)
(35000, 13) (15000, 13) (35000,) (15000,)

```

Confusion Matrix for Train data



Confusion Matrix for Test data



AUC (Train): 0.9824382772456348

AUC (Test): 0.9679149166040206

F1 SCORE (Train) : 0.9711344670786349

F1 SCORE (Test) : 0.9600679694137637

RECALL (Train): 0.9864645370871684

RECALL (Test): 0.9847884645856441

PRECISION (Train) : 0.9562735771690996

PRECISION (Test) : 0.9365581675708258

```
In [42]: merged_features_vectorizer = mapper.features[0][1].get_feature_names() + mapper.features[1][1].get_feature_names()
         important_features(merged_features_vectorizer, nb_optimal, 10)
```

Important words in negative reviews

	Class Label 0
466.440	not
306.290	not
134.836	like
126.875	product
121.272	would
115.669	taste
109.655	taste
94.308	disappointed
92.847	one
86.631	like

Important words in positive reviews

	Class Label 1
1496.994	great
1066.364	good
907.671	not
845.669	love
845.563	best
803.768	delicious
661.154	tea
644.066	great
621.736	tea

614.086 good

As we can see that by consider 'review text', 'summary text' and 'number of words' in review text all together, AUC value is changed from 0.9560 to 0.96933

8 [6] Conclusions

```
In [63]: import pandas as pd
         from prettytable import PrettyTable

         print(color.BOLD+'\t\t\t\t Naive Bayes \t'+color.END)
         print('\n')

         print(color.BOLD+'For BOW and TFIDF, We have considered 100k points'+color.END)
         print(color.BOLD+'For BOW- Additional Feature and TFIDF- Additional Feature, We have c

         x = PrettyTable()
         x.field_names = ['Metric','BOW','BOW-Additional Feature', 'TFIDF', 'TFIDF- Additional

         x.add_row(["Alpha Value ", 0.5,0.5,0.1,0.1])

         x.add_row(["AUC Train ", 0.97446,0.98161,0.98461,0.98243])
         x.add_row(["AUC Test ", 0.94502,0.95665,0.9560,0.96791])

         x.add_row(["F1 SCORE Train ", 0.96266,0.96842,0.96619,0.97113])
         x.add_row(["F1 SCORE Test ", 0.94953,0.95462,0.949638,0.96006])

         x.add_row(["RECALL Train ",0.94723,0.95276,0.99048,0.9864])
         x.add_row(["RECALL Test ", 0.97729,0.942481,0.9870,0.98478])

         x.add_row(["PRECISION Train ", 0.97861,0.98461,0.94306,0.95627])
         x.add_row(["PRECISION Test ",0.96210,0.96707,0.91498,0.93655])

         print('\n')
         print(x)
```

Naive Bayes

For BOW and TFIDF, We have considered 100k points

For BOW- Additional Feature and TFIDF- Additional Feature, We have considered 50k points

```
+-----+-----+-----+-----+-----+
|      Metric      |  BOW  | BOW-Additional Feature |  TFIDF  | TFIDF- Additional Features
```

+-----+		+-----+		+-----+	
	Alpha Value	0.5	0.5	0.1	0.1
	AUC Train	0.97446	0.98161	0.98461	0.98243
	AUC Test	0.94502	0.95665	0.956	0.96791
	F1 SCORE Train	0.96266	0.96842	0.96619	0.97113
	F1 SCORE Test	0.94953	0.95462	0.949638	0.96006
	RECALL Train	0.94723	0.95276	0.99048	0.9864
	RECALL Test	0.97729	0.942481	0.987	0.98478
	PRECISION Train	0.97861	0.98461	0.94306	0.95627
	PRECISION Test	0.9621	0.96707	0.91498	0.93655
+-----+		+-----+		+-----+	