

09 Amazon Fine Food Reviews Analysis_GDBT

June 21, 2019

1 Amazon Fine Food Reviews Analysis

Data Source: <https://www.kaggle.com/snap/amazon-fine-food-reviews>

EDA: <https://nycdatascience.com/blog/student-works/amazon-fine-foods-visualization/>

The Amazon Fine Food Reviews dataset consists of reviews of fine foods from Amazon.

Number of reviews: 568,454 Number of users: 256,059 Number of products: 74,258 Timespan: Oct 1999 - Oct 2012 Number of Attributes/Columns in data: 10

Attribute Information:

1. Id
2. ProductId - unique identifier for the product
3. UserId - unique identifier for the user
4. ProfileName
5. HelpfulnessNumerator - number of users who found the review helpful
6. HelpfulnessDenominator - number of users who indicated whether they found the review helpful or not
7. Score - rating between 1 and 5
8. Time - timestamp for the review
9. Summary - brief summary of the review
10. Text - text of the review

Objective: Given a review, determine whether the review is positive (rating of 4 or 5) or negative (rating of 1 or 2).

[Q] How to determine if a review is positive or negative? [Ans] We could use Score/Rating. A rating of 4 or 5 can be considered as a positive review. A rating of 1 or 2 can be considered as negative one. A review of rating 3 is considered neutral and such reviews are ignored from our analysis. This is an approximate and proxy way of determining the polarity (positivity/negativity) of a review.

2 [1]. Reading Data

2.1 [1.1] Loading the data

The dataset is available in two forms 1. .csv file 2. SQLite Database

In order to load the data, We have used the SQLITE dataset as it is easier to query the data and visualise the data efficiently.

Here as we only want to get the global sentiment of the recommendations (positive or negative), we will purposefully ignore all Scores equal to 3. If the score is above 3, then the recommendation will be set to "positive". Otherwise, it will be set to "negative".

```
In [1]: %matplotlib inline
import warnings
warnings.filterwarnings("ignore")

import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer

import re
# Tutorial about Python regular expressions: https://pymotw.com/2/re/
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle

from tqdm import tqdm
import os

# importing Cross validation libs
from sklearn.model_selection import train_test_split
from sklearn.model_selection import cross_val_score
from sklearn import model_selection

# Python script for confusion matrix creation.
from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score
```

```

from sklearn.metrics import classification_report
# ROC , AUC curve
# roc curve and auc
from sklearn.datasets import make_classification
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score
from matplotlib import pyplot
from sklearn.metrics import roc_curve, auc
# kFold
from sklearn.model_selection import KFold

from sklearn.model_selection import GridSearchCV

import seaborn as sns

from sklearn.model_selection import TimeSeriesSplit
from sklearn.model_selection import RandomizedSearchCV

from joblib import dump, load
from sklearn_pandas import DataFrameMapper
from sklearn.metrics import f1_score, recall_score, precision_score
import xgboost as xgb
from xgboost.sklearn import XGBClassifier

from wordcloud import WordCloud

```

```

In [2]: # using SQLite Table to read data.
con = sqlite3.connect(r'/home/pranay/ML datasource/amazon-fine-food-reviews/database.s

# filtering only positive and negative reviews i.e.
# not taking into consideration those reviews with Score=3
# SELECT * FROM Reviews WHERE Score != 3 LIMIT 500000, will give top 500000 data points
# you can change the number to any other number based on your computing power

# filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 LIMIT 5
# for tsne assignment you can take 5k data points

filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 LIMIT 8500

# Give reviews with Score>3 a positive rating(1), and reviews with a score<3 a negativ
def partition(x):
    if x < 3:
        return 0
    return 1

#changing reviews with score less than 3 to be positive and vice-versa
actualScore = filtered_data['Score']
positiveNegative = actualScore.map(partition)

```

```

filtered_data['Score'] = positiveNegative
print("Number of data points in our data", filtered_data.shape)
filtered_data.head(3)

```

Number of data points in our data (85000, 10)

```

Out[2]:
   Id  ProductId      UserId      ProfileName \
0   1  B001E4KFG0  A3SGXH7AUHU8GW      delmartian
1   2  B00813GRG4  A1D87F6ZCVE5NK      dll pa
2   3  B000LQOCHO  ABXLMWJIXXAIN  Natalia Corres "Natalia Corres"

   HelpfulnessNumerator  HelpfulnessDenominator  Score      Time \
0                      1                      1      1  1303862400
1                      0                      0      0  1346976000
2                      1                      1      1  1219017600

           Summary      Text
0  Good Quality Dog Food  I have bought several of the Vitality canned d...
1    Not as Advertised  Product arrived labeled as Jumbo Salted Peanut...
2  "Delight" says it all  This is a confection that has been around a fe...

```

```

In [3]: display = pd.read_sql_query("""
SELECT UserId, ProductId, ProfileName, Time, Score, Text, COUNT(*)
FROM Reviews
GROUP BY UserId
HAVING COUNT(*)>1
""", con)

```

```

In [4]: print(display.shape)
display.head()

```

(80668, 7)

```

Out[4]:
   UserId      ProductId      ProfileName      Time  Score \
0 #oc-R115TNMSPFT9I7  B005ZBZLT4      Breyton  1331510400      2
1 #oc-R11D9D7SHXIJB9  B005HG9ESG  Louis E. Emory "hoppy"  1342396800      5
2 #oc-R11DNU2NBKQ23Z  B005ZBZLT4      Kim Cieszykowski  1348531200      1
3 #oc-R1105J5ZVQE25C  B005HG9ESG      Penguin Chick  1346889600      5
4 #oc-R12KPBODL2B5ZD  B007OSBEV0  Christopher P. Presta  1348617600      1

           Text  COUNT(*)
0  Overall its just OK when considering the price...      2
1  My wife has recurring extreme muscle spasms, u...      3
2  This coffee is horrible and unfortunately not ...      2
3  This will be the bottle that you grab from the...      3
4  I didnt like this coffee. Instead of telling y...      2

```

```

In [5]: display[display['UserId']=='AZY10LLTJ71NX']

```

```
Out [5]:
```

	UserId	ProductId	ProfileName	Time	\
80638	AZY10LLTJ71NX	B001ATMQK2	undertheshrine	"undertheshrine"	1296691200

	Score	Text	COUNT(*)
80638	5	I bought this 6 pack because for the price tha...	5

```
In [6]: display['COUNT(*)'].sum()
```

```
Out [6]: 393063
```

3 [2] Exploratory Data Analysis

3.1 [2.1] Data Cleaning: Deduplication

It is observed (as shown in the table below) that the reviews data had many duplicate entries. Hence it was necessary to remove duplicates in order to get unbiased results for the analysis of the data. Following is an example:

```
In [7]: display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND UserId="AR5J8UI46CURR"
ORDER BY ProductID
""", con)
display.head()
```

```
Out [7]:
```

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	\
0	78445	B000HDL1RQ	AR5J8UI46CURR	Geetha Krishnan	2	
1	138317	B000HDOPYC	AR5J8UI46CURR	Geetha Krishnan	2	
2	138277	B000HDOPYM	AR5J8UI46CURR	Geetha Krishnan	2	
3	73791	B000HDOPZG	AR5J8UI46CURR	Geetha Krishnan	2	
4	155049	B000PAQ75C	AR5J8UI46CURR	Geetha Krishnan	2	

	HelpfulnessDenominator	Score	Time	\
0	2	5	1199577600	
1	2	5	1199577600	
2	2	5	1199577600	
3	2	5	1199577600	
4	2	5	1199577600	

	Summary	\
0	LOACKER QUADRATINI VANILLA WAFERS	
1	LOACKER QUADRATINI VANILLA WAFERS	
2	LOACKER QUADRATINI VANILLA WAFERS	
3	LOACKER QUADRATINI VANILLA WAFERS	
4	LOACKER QUADRATINI VANILLA WAFERS	

Text

```

0 DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
1 DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
2 DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
3 DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
4 DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...

```

As it can be seen above that same user has multiple reviews with same values for HelpfulnessNumerator, HelpfulnessDenominator, Score, Time, Summary and Text and on doing analysis it was found that ProductId=B000HDOPZG was Loacker Quadratini Vanilla Wafer Cookies, 8.82-Ounce Packages (Pack of 8) ProductId=B000HDL1RQ was Loacker Quadratini Lemon Wafer Cookies, 8.82-Ounce Packages (Pack of 8) and so on

It was inferred after analysis that reviews with same parameters other than ProductId belonged to the same product just having different flavour or quantity. Hence in order to reduce redundancy it was decided to eliminate the rows having same parameters.

The method used for the same was that we first sort the data according to ProductId and then just keep the first similar product review and delete the others. for eg. in the above just the review for ProductId=B000HDL1RQ remains. This method ensures that there is only one representative for each product and deduplication without sorting would lead to possibility of different representatives still existing for the same product.

```

In [8]: #Sorting data according to ProductId in ascending order
sorted_data=filtered_data.sort_values('ProductId', axis=0, ascending=True, inplace=False)

```

```

In [9]: #Deduplication of entries
final=sorted_data.drop_duplicates(subset={"UserId","ProfileName","Time","Text"}, keep='first')
final.shape

```

```

Out[9]: (75844, 10)

```

```

In [10]: #Checking to see how much % of data still remains
(final['Id'].size*1.0)/(filtered_data['Id'].size*1.0)*100

```

```

Out[10]: 89.22823529411764

```

Observation:- It was also seen that in two rows given below the value of HelpfulnessNumerator is greater than HelpfulnessDenominator which is not practically possible hence these two rows too are removed from calculations

```

In [11]: display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND Id=44737 OR Id=64422
ORDER BY ProductID
""", con)

display.head()

```

```

Out[11]:
   Id  ProductId  UserId  ProfileName \
0  64422  B000MIDR0Q  A161DK06JJMCF  J. E. Stephens "Jeanne"

```

	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	\
0	3	1	5	1224892800	
1	3	2	4	1212883200	

	Summary	\
0	Bought This for My Son at College	
1	Pure cocoa taste with crunchy almonds inside	

	Text
0	My son loves spaghetti so I didn't hesitate or...
1	It was almost a 'love at first bite' - the per...

```
In [12]: final=final[final.HelpfulnessNumerator<=final.HelpfulnessDenominator]
```

```
In [13]: #Before starting the next phase of preprocessing lets see the number of entries left
print(final.shape)
```

```
#How many positive and negative reviews are present in our dataset?
final['Score'].value_counts()
```

```
(75842, 10)
```

```
Out[13]: 1    63459
         0    12383
         Name: Score, dtype: int64
```

4 [3] Preprocessing

4.1 [3.1]. Preprocessing Review Text

Now that we have finished deduplication our data requires some preprocessing before we go on further with analysis and making the prediction model.

Hence in the Preprocessing phase we do the following in the order below:-

1. Begin by removing the html tags
2. Remove any punctuations or limited set of special characters like , or . or # etc.
3. Check if the word is made up of english letters and is not alpha-numeric
4. Check to see if the length of the word is greater than 2 (as it was researched that there is no adjective in 2-letters)
5. Convert the word to lowercase
6. Remove Stopwords
7. Finally Snowball Stemming the word (it was observed to be better than Porter Stemming)

After which we collect the words used to describe positive and negative reviews

```
In [14]: # printing some random reviews
sent_0 = final['Text'].values[0]
print(sent_0)
print("="*50)

sent_1000 = final['Text'].values[1000]
print(sent_1000)
print("="*50)

sent_1500 = final['Text'].values[1500]
print(sent_1500)
print("="*50)

sent_4900 = final['Text'].values[4900]
print(sent_4900)
print("="*50)
```

My dogs loves this chicken but its a product from China, so we wont be buying it anymore. Its
=====
These were just adorable when used on a bee hive cake! In the right setting they add that spee
=====
This is the BEST!
Years ago, I used it & loved it! Moved & could not find it again. I hav
=====
This product arrived in a timely manner, in good condition, and it was a hit with the family w
=====

```
In [15]: # remove urls from text python: https://stackoverflow.com/a/40823105/4084039
sent_0 = re.sub(r"http\S+", "", sent_0)
sent_1000 = re.sub(r"http\S+", "", sent_1000)
sent_150 = re.sub(r"http\S+", "", sent_1500)
sent_4900 = re.sub(r"http\S+", "", sent_4900)

print(sent_0)
```

My dogs loves this chicken but its a product from China, so we wont be buying it anymore. Its

```
In [16]: # https://stackoverflow.com/questions/16206380/python-beautifulsoup-how-to-remove-all
from bs4 import BeautifulSoup

soup = BeautifulSoup(sent_0, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_1000, 'lxml')
text = soup.get_text()
print(text)
```



```

print("="*50)

soup = BeautifulSoup(sent_1500, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_4900, 'lxml')
text = soup.get_text()
print(text)

```

My dogs loves this chicken but its a product from China, so we wont be buying it anymore. Its
=====

These were just adorable when used on a bee hive cake! In the right setting they add that spe
=====

This is the BEST!Years ago, I used it & loved it! Moved & could not find it again. I have purch
=====

This product arrived in a timely manner, in good condition, and it was a hit with the family w

In [17]: # <https://stackoverflow.com/a/47091490/4084039>

```

import re

def decontracted(phrase):
    # specific
    phrase = re.sub(r"won't", "will not", phrase)
    phrase = re.sub(r"can't", "can not", phrase)

    # general
    phrase = re.sub(r"n't", " not", phrase)
    phrase = re.sub(r"\ 're", " are", phrase)
    phrase = re.sub(r"\ 's", " is", phrase)
    phrase = re.sub(r"\ 'd", " would", phrase)
    phrase = re.sub(r"\ 'll", " will", phrase)
    phrase = re.sub(r"\ 't", " not", phrase)
    phrase = re.sub(r"\ 've", " have", phrase)
    phrase = re.sub(r"\ 'm", " am", phrase)
    return phrase

```

```

In [18]: sent_1500 = decontracted(sent_1500)
print(sent_1500)
print("="*50)

```

This is the BEST!
Years ago, I used it & loved it! Moved & could not find it again. I have
=====

In [19]: #remove words with numbers python: <https://stackoverflow.com/a/18082370/4084039>

```

sent_0 = re.sub("\S*\d\S*", "", sent_0).strip()
print(sent_0)

```

My dogs loves this chicken but its a product from China, so we wont be buying it anymore. Its

```
In [20]: #remove spacial character: https://stackoverflow.com/a/5843547/4084039
sent_1500 = re.sub('[^A-Za-z0-9]+', ' ', sent_1500)
print(sent_1500)
```

This is the BEST br Years ago I used it loved it Moved could not find it again I have purchased

```
In [21]: # https://gist.github.com/sebleier/554280
# we are removing the words from the stop words list: 'no', 'nor', 'not'
# <br /><br /> ==> after the above steps, we are getting "br br"
# we are including them into stop words list
# instead of <br /> if we have <br/> these tags would have reuvmoved in the 1st step

stopwords= set(['br', 'the', 'i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves',
                "you'll", "you'd", 'your', 'yours', 'yourself', 'yourselves', 'he', 'him',
                'she', "she's", 'her', 'hers', 'herself', 'it', "it's", 'its', 'itself',
                'theirs', 'themselves', 'what', 'which', 'who', 'whom', 'this', 'that', "that's",
                'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'have', 'has', 'had',
                'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or', 'because', 'as', 'until',
                'at', 'by', 'for', 'with', 'about', 'against', 'between', 'into', 'through',
                'above', 'below', 'to', 'from', 'up', 'down', 'in', 'out', 'on', 'off', 'over',
                'then', 'once', 'here', 'there', 'when', 'where', 'why', 'how', 'all', 'any',
                'most', 'other', 'some', 'such', 'only', 'own', 'same', 'so', 'than', 'too', 'very',
                's', 't', 'can', 'will', 'just', 'don', "don't", 'should', "should've", 'n',
                've', 'y', 'ain', 'aren', "aren't", 'couldn', "couldn't", 'didn', "didn't",
                "hadn't", 'hasn', "hasn't", 'haven', "haven't", 'isn', "isn't", 'ma', 'mi',
                "mustn't", 'needn', "needn't", 'shan', "shan't", 'shouldn', "shouldn't",
                'won', "won't", 'wouldn', "wouldn't"])
```

```
In [22]: #filtered out whole reviews
```

```
from bs4 import BeautifulSoup
# Combining all the above stundents
from tqdm import tqdm
# tqdm is for printing the status bar
word_counter = []
def filterised_text(text):
    preprocessed_text = []
    for sentence in tqdm(text):
        sentence = re.sub(r"http\S+", "", sentence)
        sentence = BeautifulSoup(sentence, 'lxml').get_text()
        sentence = decontracted(sentence)
        sentence = re.sub("\S*\d\S*", "", sentence).strip()
        sentence = re.sub('[^A-Za-z]+', ' ', sentence)
        # https://gist.github.com/sebleier/554280
        sentence = ' '.join(e.lower() for e in sentence.split() if e.lower() not in stopwords)
```

```

        count = len(sentence.split())
        word_counter.append(count)
        preprocessed_text.append(sentence.strip())
    return preprocessed_text

```

```

In [23]: preprocessed_reviews = filterised_text(final['Text'].values)
        final['preprocessed_reviews'] = preprocessed_reviews
        preprocessed_reviews[1822]

```

100%|| 75842/75842 [00:25<00:00, 2930.99it/s]

Out[23]: 'fell love product england boyfriend sandwich branston never want eat sandwich without'

```

In [24]: final['numbers_of_words'] = word_counter
        word_counter[1822]

```

Out[24]: 34

4.1.1 [3.2] Preprocessing Review Summary

```

In [25]: preprocessed_summary = filterised_text(final['Summary'].values)
        final['preprocessed_summary'] = preprocessed_summary
        preprocessed_summary[1822]

```

100%|| 75842/75842 [00:14<00:00, 5368.92it/s]

Out[25]: 'branstolicious'

```

In [26]: avg_w2v_trained_model_100000 = '/home/pranay/ML trained models/W2V/avg_w2v_trained_model_100000'
        avg_w2v_test_model_100000 = '/home/pranay/ML trained models/W2V/avg_w2v_test_model_100000'

```

```

        w2v_tf_idf_trained_model_100000 = '/home/pranay/ML trained models/W2V_TFIDF/w2v_tf_idf_trained_model_100000'
        w2v_tf_idf_test_model_100000 = '/home/pranay/ML trained models/W2V_TFIDF/w2v_tf_idf_test_model_100000'

```

```

In [37]: depth_ = [2,5,7,10,25]
        depth_ = np.asarray(depth_)

```

```

        estimators = [25,50,100,500,1000]
        estimators_list = np.asarray(estimators)

```

```

def finding_best_hyperparam(X_tr,y_tr):
    # instantiate a GBDT model
    xgb = XGBClassifier(class_weight='balanced', random_state=1)

    param_grid=dict(n_estimators=estimators_list,max_depth=depth_)

```

```

#For time based splitting
tscv = TimeSeriesSplit(n_splits=10)

# instantiate the training grid search model
train_grid = GridSearchCV(xgb, param_grid, cv=tscv, scoring='roc_auc',n_jobs=-1,

# fit the training data to train model
train_grid.fit(X_tr, y_tr)

return train_grid

# https://machinelearningmastery.com/roc-curves-and-precision-recall-curves-for-class

# plot AUC curve
def plotAUC_ROC(model,X_train, y_train, X_test, y_test):
    # predict probabilities
    test_probs = model.predict_proba(X_test)
    train_probs = model.predict_proba(X_train)

    # keep probabilities for the positive outcome only
    test_probs = test_probs[:, 1]
    train_probs = train_probs[:, 1]

    # calculate AUC
    test_auc = roc_auc_score(y_test, test_probs)
    train_auc = roc_auc_score(y_train, train_probs)

    # calculate roc curve
    train_fpr, train_tpr, thresholds = roc_curve(y_train, train_probs)
    test_fpr, test_tpr, thresholds2 = roc_curve(y_test, test_probs)

    # plot no skill
    pyplot.plot([0, 1], [0, 1], linestyle='--')

    # plot the roc curve for the model
    pyplot.plot(train_fpr, train_tpr, 'r',marker='.', label="train AUC =" +str(train_auc))
    pyplot.plot(test_fpr, test_tpr, 'b',marker='.',label="test AUC =" +str(test_auc))

    pyplot.legend()
    pyplot.xlabel("K: hyperparameter")
    pyplot.ylabel("AUC")
    pyplot.title("ERROR PLOTS")
    # show the plot
    pyplot.show()

return train_auc, test_auc

```

```

# https://www.geeksforgeeks.org/confusion-matrix-machine-learning/
def plotConfusionMatrix(y_test, pred):
    # calculate confusion matrix
    cm = confusion_matrix(y_test, pred)
    class_label = ['negative', 'positive']
    df_conf_matrix = pd.DataFrame(cm, index=class_label, columns=class_label)
    # heatmap --> Plot rectangular data as a color-encoded matrix.
    sns.heatmap(df_conf_matrix, annot=True, fmt='d')
    # give title to graph
    plt.title("Confusion Matrix")
    # mention axis label
    plt.xlabel("Predicted")
    plt.ylabel("Actual")
    # show the plot
    plt.show()

class color:
    PURPLE = '\033[95m'
    CYAN = '\033[96m'
    DARKCYAN = '\033[36m'
    BLUE = '\033[94m'
    GREEN = '\033[92m'
    YELLOW = '\033[93m'
    RED = '\033[91m'
    BOLD = '\033[1m'
    UNDERLINE = '\033[4m'
    END = '\033[0m'

# https://qiita.com/bmj0114/items/8009f282c99b77780563
def plotHeatMap(trained_model, param):
    if param == 'trained':
        scores = trained_model.cv_results_['mean_train_score'].reshape(len(estimators_list))
    else:
        scores = trained_model.cv_results_['mean_test_score'].reshape(len(estimators_list))

    plt.figure(figsize=(16, 12))
    sns.heatmap(scores, annot=True, cmap=plt.cm.hot, fmt=".3f", xticklabels=estimators_list)
    plt.xlabel('n_estimators')
    plt.ylabel('max_depth')
    plt.xticks(np.arange(len(estimators_list)), estimators_list)
    plt.yticks(np.arange(len(depth_)), depth_)
    plt.title('Grid Search AUC Score')
    plt.show()

```

5 [4] Featurization

5.0.1 Splitting data

We have considered 85 k points

```
In [78]: X = final['preprocessed_reviews']
         y = final['Score']

         # split the data set into train and test
         X_train, x_test, y_train, y_test = model_selection.train_test_split(X, y, test_size=0.25)

         print(X_train.shape, x_test.shape, y_train.shape, y_test.shape)

(53089,) (22753,) (53089,) (22753,)
```

5.1 [4.1] BAG OF WORDS

```
In [79]: ##BoW

         count_vect = CountVectorizer(ngram_range=(1,2), min_df=10) #in scikit-learn

         # train data
         X_train_bow = count_vect.fit_transform(X_train)

         # test data
         x_test_bow = count_vect.transform(x_test)

         print('X_train_bow',X_train_bow.shape)
         print('==='*10)
         print('x_test_bow',x_test_bow.shape)

X_train_bow (53089, 31683)
=====
x_test_bow (22753, 31683)
```

5.1.1 Hyper param Tuning using GridSearch

finding 'max depth' & 'esimate models' which have maximum AUC Score

```
In [31]: bow_train_path = '/home/pranay/ML Hyperparam Tune/GBDT/bow_train_hyperparam_tuned'
         exists = os.path.isfile(bow_train_path)

         if exists:
             print("yes exists")
             bow_train = load(bow_train_path)
         else:
```

```

print("not exists")
bow_train = finding_best_hyperparam(X_train_bow,y_train)
dump(bow_train,bow_train_path )

# view the complete results (list of named tuples)
print("=====Training=====")
print (bow_train.best_score_)
print (bow_train.best_params_)
print (bow_train.best_estimator_)

# plotAccuracyGraph(bow_train,'n_estimators')
best_estimators = bow_train.best_params_.get("n_estimators","")
best_depth_size = bow_train.best_params_.get("max_depth", "")

best_estimators, best_depth_size


yes exists
=====Training=====
0.9377287074476324
{'max_depth': 5, 'n_estimators': 1000}
XGBClassifier(base_score=0.5, booster='gbtree', class_weight='balanced',
               colsample_bylevel=1, colsample_bytree=1, gamma=0, learning_rate=0.1,
               max_delta_step=0, max_depth=5, min_child_weight=1, missing=nan,
               n_estimators=1000, n_jobs=1, nthread=None,
               objective='binary:logistic', random_state=1, reg_alpha=0,
               reg_lambda=1, scale_pos_weight=1, seed=None, silent=True,
               subsample=1)

Out[31]: (1000, 5)

In [32]: print('\n'+color.BOLD +'AUC Train data' +color.END)
         plotHeatMap(bow_train,'trained')


```

AUC Train data



```
In [33]: print('\n'+color.BOLD + 'AUC Validation data'+color.END)
          plotHeatMap(bow_train, 'test')
```

AUC Validation data



5.1.2 Applying GBDT on BOW

```
In [81]: optimal_model = XGBClassifier(base_score=0.5, booster='gbtree', class_weight='balanced',
    colsample_bylevel=1, colsample_bytree=1, gamma=0, learning_rate=0.1,
    max_delta_step=0, max_depth=4, min_child_weight=1, missing=None,
    n_estimators=100, n_jobs=1, nthread=None,
    objective='binary:logistic', random_state=1, reg_alpha=0,
    reg_lambda=1, scale_pos_weight=1, seed=None, silent=True,
    subsample=1)
```

```
# fitting the model
```

```
optimal_model.fit(X_train_bow, y_train)
```

```
# predict the response
```

```
test_pred = optimal_model.predict(x_test_bow)
```

```
train_pred = optimal_model.predict(X_train_bow)
```

```

print('\n'+color.RED+'Max Depth : '+color.END+color.BOLD+str(4)+color.END)
print('\n'+color.RED+'Best Estimator : '+color.END+color.BOLD+str(100)+color.END)

# plot confusion matrix

print('\n'+color.BOLD +'Confusion Matrix for Train data'+color.END)
plotConfusionMatrix(y_train,train_pred)

print('\n'+color.BOLD +'Confusion Matrix for Test data'+color.END)
plotConfusionMatrix(y_test,test_pred)

# plot AUC
train_auc,test_auc = plotAUC_ROC(optimal_model,X_train_bow, y_train,x_test_bow, y_test)
print('\n'+color.RED+'AUC (Train): '+color.END+color.BOLD+str(train_auc)+color.END)
print('\n'+color.RED+'AUC (Test): '+color.END+color.BOLD+str(test_auc)+color.END)

# f1 score
score = f1_score(y_test,test_pred)
print('\n'+color.RED+'F1 SCORE (Train) : '+color.END+color.BOLD+str(f1_score(y_train,train_pred))+color.END)
print('\n'+color.RED+'F1 SCORE (Test) : '+color.END+color.BOLD+str(score)+color.END)

# recall
recall = metrics.recall_score(y_test, test_pred)
print('\n'+color.RED+'RECALL (Train): '+color.END+color.BOLD+str(metrics.recall_score(y_train,train_pred))+color.END)
print('\n'+color.RED+'RECALL (Test): '+color.END+color.BOLD+str(recall)+color.END)


# precision
precision = metrics.precision_score(y_test, test_pred)
print('\n'+color.RED+'PRECISION (Train) : '+color.END+color.BOLD+str(metrics.precision_score(y_train,train_pred))+color.END)
print('\n'+color.RED+'PRECISION (Test) : '+color.END+color.BOLD+str(precision)+color.END)

```


Max Depth : 4


Best Estimator : 100

Confusion Matrix for Train data



Confusion Matrix for Test data





AUC (Train): 0.9148762168853501

AUC (Test): 0.9016486570976912

F1 SCORE (Train) : 0.9360325598604579

F1 SCORE (Test) : 0.9311800874138826

RECALL (Train): 0.9955475601529121

RECALL (Test): 0.9929432829532887

PRECISION (Train) : 0.8832319201995013

PRECISION (Test) : 0.8766505486330668

5.1.3 Top 20 features

```
In [82]: topn_class = sorted(zip(optimal_model.feature_importances_, count_vect.get_feature_names()))
         top_words_ = ''
```

```

    for feature, value in topn_class:
        print(feature, value)
        top_words_ += ' ' +value

0.017997468 waste money
0.017685018 return
0.014468065 delicious
0.014412325 terrible
0.014196279 great
0.013957699 worst
0.013775373 awful
0.013386117 not disappointed
0.013293953 money
0.01315962 refund
0.013138805 disappointed
0.012694619 not
0.012359969 threw
0.012193472 not good
0.012154945 horrible
0.011763318 favorite
0.011446781 not buy
0.0113734165 nice
0.011363316 bad
0.011279832 not recommend

```

In [83]: # <https://www.geeksforgeeks.org/generating-word-cloud-python/>

```

wordcloud = WordCloud(width = 800, height = 800,
                       background_color = 'white',
                       min_font_size = 10).generate(top_words_)

# plot the WordCloud image
plt.figure(figsize = (8, 8), facecolor = None)
plt.imshow(wordcloud)
plt.axis("off")
plt.tight_layout(pad = 0)

plt.show()

```

nice waste
good delicious
favorite threw
great return
money awful
refund
buybad terrible
disappointed
horrible worse
recommend

5.2 Feature Engineering

Till now we only consider Text review as feature, we are adding some extra feature like **review summary** and **number of words** in review and test our model improves efficiency or not.

We have considered on 30000 points due to memory issue.

In [32]: # <https://sondosatwi.wordpress.com/2017/08/01/using-text-data-and-dataframemapper-in->

```
X = final[:30000]
y = final['Score'][:30000]

# split the data set into train and test
X_train, x_test, y_train, y_test = model_selection.train_test_split(X, y, test_size=0
```

```

print(X_train.shape, x_test.shape, y_train.shape, y_test.shape)

mapper = DataFrameMapper([
    ('preprocessed_reviews', CountVectorizer(ngram_range=(1,3), min_df=10)),
    ('preprocessed_summary', CountVectorizer(ngram_range=(1,3), min_df=10)),
    ('numbers_of_words', None),
])

train_features = mapper.fit_transform(X_train)
test_features = mapper.transform(x_test)

optimal_model = XGBClassifier(base_score=0.5, booster='gbtree', class_weight='balanced',
    colsample_bylevel=1, colsample_bytree=1, gamma=0, learning_rate=0.1,
    max_delta_step=0, max_depth=4, min_child_weight=1,
    n_estimators=100, n_jobs=-1, nthread=None,
    objective='binary:logistic', random_state=1, reg_alpha=0,
    reg_lambda=1, scale_pos_weight=1, seed=None, silent=True,
    subsample=1)

# fitting the model
optimal_model.fit(train_features,y_train)

# predict the response
test_pred = optimal_model.predict(test_features)
train_pred = optimal_model.predict(train_features)

# plot confusion matrix

print('\n'+color.BOLD + 'Confusion Matrix for Train data'+color.END)
plotConfusionMatrix(y_train,train_pred)

print('\n'+color.BOLD + 'Confusion Matrix for Test data'+color.END)
plotConfusionMatrix(y_test,test_pred)

# plot AUC
train_auc,test_auc = plotAUC_ROC(optimal_model,train_features, y_train,test_features,
print('\n'+color.RED+'AUC (Train): '+color.END+color.BOLD+str(train_auc)+color.END)
print('\n'+color.RED+'AUC (Test): '+color.END+color.BOLD+str(test_auc)+color.END)

# f1 score
score = f1_score(y_test,test_pred)
print('\n'+color.RED+'F1 SCORE (Train) : '+color.END+color.BOLD+str(f1_score(y_train,train_pred))+color.END)
print('\n'+color.RED+'F1 SCORE (Test) : '+color.END+color.BOLD+str(score)+color.END)

# recall
recall = metrics.recall_score(y_test, test_pred)
print('\n'+color.RED+'RECALL (Train): '+color.END+color.BOLD+str(metrics.recall_score(y_train,train_pred))+color.END)
print('\n'+color.RED+'RECALL (Test): '+color.END+color.BOLD+str(recall)+color.END)

```


```

# precision
precision = metrics.precision_score(y_test, test_pred)
print('\n'+color.RED+'PRECISION (Train) : '+color.END+color.BOLD+str(metrics.precision))
print('\n'+color.RED+'PRECISION (Test) : '+color.END+color.BOLD+str(precision)+color.END)



(21000, 13) (9000, 13) (21000,) (9000,)

```

Confusion Matrix for Train data



Confusion Matrix for Test data



```

AUC (Train): 0.9945409933070557

AUC (Test): 0.9670874541400183

F1 SCORE (Train) : 0.9822161422708618

F1 SCORE (Test) : 0.9606828763579928

RECALL (Train): 0.9957545567757274

RECALL (Test): 0.983450284655104

PRECISION (Train) : 0.96904092987385

PRECISION (Test) : 0.9389457717102768

```

5.3 [4.3] TF-IDF

```

In [41]: X = final['preprocessed_reviews']
        y = final['Score']

        # split the data set into train and test
        X_train, x_test, y_train, y_test = model_selection.train_test_split(X, y, test_size=0.2)

        print(X_train.shape, x_test.shape, y_train.shape, y_test.shape)

        tf_idf_vect = TfidfVectorizer(ngram_range=(1,3), min_df=10) #in scikit-learn

        # train data
        X_train_tfidf = tf_idf_vect.fit_transform(X_train)

        # test data
        x_test_tfidf = tf_idf_vect.transform(x_test)

        print('X_train_tfidf', X_train_tfidf.shape)
        print('===='*10)
        print('x_test_tfidf', x_test_tfidf.shape)

(53089,) (22753,) (53089,) (22753,)
X_train_tfidf (53089, 34425)
=====
x_test_tfidf (22753, 34425)

```

5.3.1 Hyper param Tuning using GridSearch

finding 'max depth' & 'esimate models' which have maximum AUC Score

```
In [42]: tfidf_train_path = '/home/pranay/ML Hyperparam Tune/GBDT/tfidf_train_hyperparam_tuned
exists = os.path.isfile(tfidf_train_path)

if exists:
    print("yes exists")
    tfidf_train = load(tfidf_train_path)
else:
    print("not exists")
    tfidf_train = finding_best_hyperparam(X_train_tfidf,y_train)
    dump(tfidf_train,tfidf_train_path )


# view the complete results (list of named tuples)
print("=====Training=====")
print (tfidf_train.best_score_)
print (tfidf_train.best_params_)
print (tfidf_train.best_estimator_)

best_depth_size = tfidf_train.best_params_.get("max_depth", "")
best_estimators = tfidf_train.best_params_.get("n_estimators", "")

yes exists
=====Training=====
0.934506456797819
{'max_depth': 5, 'n_estimators': 1000}
XGBClassifier(base_score=0.5, booster='gbtree', class_weight='balanced',
              colsample_bylevel=1, colsample_bytree=1, gamma=0, learning_rate=0.1,
              max_delta_step=0, max_depth=5, min_child_weight=1, missing=nan,
              n_estimators=1000, n_jobs=1, nthread=None,
              objective='binary:logistic', random_state=1, reg_alpha=0,
              reg_lambda=1, scale_pos_weight=1, seed=None, silent=True,
              subsample=1)


In [43]: print('\n'+color.BOLD +'AUC Train data'+color.END)
          plotHeatMap(tfidf_train,'trained')
```

AUC Train data



```
In [44]: print('\n'+color.BOLD + 'AUC Validation data'+color.END)
          plotHeatMap(tfidf_train, 'test')
```

AUC Validation data



```
In [45]: optimal_model = XGBClassifier(base_score=0.5, booster='gbtree', class_weight='balanced',
    colsample_bylevel=1, colsample_bytree=1, gamma=0, learning_rate=0.1,
    max_delta_step=0, max_depth=3, min_child_weight=1, missing=None,
    n_estimators=100, n_jobs=-1, nthread=None,
    objective='binary:logistic', random_state=1, reg_alpha=0,
    reg_lambda=1, scale_pos_weight=1, seed=None, silent=True,
    subsample=1)

# fitting the model
optimal_model.fit(X_train_tfidf, y_train)

# predict the response
test_pred = optimal_model.predict(x_test_tfidf)
train_pred = optimal_model.predict(X_train_tfidf)

print('\n'+color.RED+'Max Depth : '+color.END+color.BOLD+str(3)+color.END)
print('\n'+color.RED+'Best Estimator : '+color.END+color.BOLD+str(100)+color.END)
```

```

# plot confusion matrix

print('\n'+color.BOLD + 'Confusion Matrix for Train data'+color.END)
plotConfusionMatrix(y_train,train_pred)

print('\n'+color.BOLD + 'Confusion Matrix for Test data'+color.END)
plotConfusionMatrix(y_test,test_pred)

# plot AUC
train_auc,test_auc = plotAUC_ROC(optimal_model,X_train_tfidf, y_train,x_test_tfidf, y_test)
print('\n'+color.RED+'AUC (Train): '+color.END+color.BOLD+str(train_auc)+color.END)
print('\n'+color.RED+'AUC (Test): '+color.END+color.BOLD+str(test_auc)+color.END)

# f1 score
score = f1_score(y_test,test_pred)
print('\n'+color.RED+'F1 SCORE (Train) : '+color.END+color.BOLD+str(f1_score(y_train,y_test))+color.END)
print('\n'+color.RED+'F1 SCORE (Test) : '+color.END+color.BOLD+str(score)+color.END)

# recall
recall = metrics.recall_score(y_test, test_pred)
print('\n'+color.RED+'RECALL (Train): '+color.END+color.BOLD+str(metrics.recall_score(y_train,y_test))+color.END)
print('\n'+color.RED+'RECALL (Test): '+color.END+color.BOLD+str(recall)+color.END)


# precision
precision = metrics.precision_score(y_test, test_pred)
print('\n'+color.RED+'PRECISION (Train) : '+color.END+color.BOLD+str(metrics.precision_score(y_train,y_test))+color.END)
print('\n'+color.RED+'PRECISION (Test) : '+color.END+color.BOLD+str(precision)+color.END)

```


Max Depth : 3


Best Estimator : 100

Confusion Matrix for Train data



Confusion Matrix for Test data





AUC (Train): 0.9028939377087551

AUC (Test): 0.8877669277067337

F1 SCORE (Train) : 0.9314343332493068

F1 SCORE (Test) : 0.9272111514675567

RECALL (Train): 0.9970766809084777

RECALL (Test): 0.9948391173837485

PRECISION (Train) : 0.8739012180219954

PRECISION (Test) : 0.8681924720805184

5.4 Top 20 Features

```
In [47]: topn_class = sorted(zip(optimal_model.feature_importances_, tf_idf_vect.get_feature_names()))
         top_words_tfidf = ''
```



```

        for feature, value in topn_class:
            print(feature, value)
            top_words_tfidf += ' ' +value

0.0234979 not
0.019388214 return
0.017660951 awful
0.017047124 horrible
0.01694386 threw
0.016529012 worst
0.015682874 money
0.015432428 delicious
0.015123611 great
0.014597757 disappointed
0.014460342 terrible
0.014275661 bad
0.01400822 not buy
0.013310471 received
0.013188662 love
0.013080963 waste money
0.012949058 best
0.012549151 perfect
0.01253835 refund
0.012405823 nice

```

In [48]: # <https://www.geeksforgeeks.org/generating-word-cloud-python/>

```

wordcloud = WordCloud(width = 800, height = 800,
                        background_color = 'white',
                        min_font_size = 10).generate(top_words_tfidf)

# plot the WordCloud image
plt.figure(figsize = (8, 8), facecolor = None)
plt.imshow(wordcloud)
plt.axis("off")
plt.tight_layout(pad = 0)

plt.show()

```



5.4.1 Feature Engineering

Till now we only consider Text review as feature, we are adding some extra feature like **review summary** and **number of words** in review and test our model improves efficiency or not.

We have considered on 30000 points due to memory issue.

In [49]: # <https://sondosatwi.wordpress.com/2017/08/01/using-text-data-and-dataframemapper-in->

```
X = final[:30000]
y = final['Score'][:30000]

# split the data set into train and test
X_train, x_test, y_train, y_test = model_selection.train_test_split(X, y, test_size=0
```

```

print(X_train.shape, x_test.shape, y_train.shape, y_test.shape)

mapper = DataFrameMapper([
    ('preprocessed_reviews', TfidfVectorizer(ngram_range=(1,3), min_df=10)),
    ('preprocessed_summary', TfidfVectorizer(ngram_range=(1,3), min_df=10)),
    ('numbers_of_words', None),
])

train_features = mapper.fit_transform(X_train)
test_features = mapper.transform(x_test)

optimal_model = XGBClassifier(base_score=0.5, booster='gbtree', class_weight='balanced',
    colsample_bylevel=1, colsample_bytree=1, gamma=0, learning_rate=0.1,
    max_delta_step=0, max_depth=3, min_child_weight=1, missing=None,
    n_estimators=100, n_jobs=-1, nthread=None,
    objective='binary:logistic', random_state=1, reg_alpha=0,
    reg_lambda=1, scale_pos_weight=1, seed=None, silent=True,
    subsample=1)

# fitting the model
optimal_model.fit(train_features,y_train)

# predict the response
test_pred = optimal_model.predict(test_features)
train_pred = optimal_model.predict(train_features)

# plot confusion matrix

print('\n'+color.BOLD + 'Confusion Matrix for Train data'+color.END)
plotConfusionMatrix(y_train,train_pred)

print('\n'+color.BOLD + 'Confusion Matrix for Test data'+color.END)
plotConfusionMatrix(y_test,test_pred)

# plot AUC
train_auc,test_auc = plotAUC_ROC(optimal_model,train_features, y_train,test_features,
print('\n'+color.RED+'AUC (Train): '+color.END+color.BOLD+str(train_auc)+color.END)
print('\n'+color.RED+'AUC (Test): '+color.END+color.BOLD+str(test_auc)+color.END)

# f1 score
score = f1_score(y_test,test_pred)
print('\n'+color.RED+'F1 SCORE (Train) : '+color.END+color.BOLD+str(f1_score(y_train,train_pred))+color.END)
print('\n'+color.RED+'F1 SCORE (Test) : '+color.END+color.BOLD+str(score)+color.END)

# recall
recall = metrics.recall_score(y_test, test_pred)
print('\n'+color.RED+'RECALL (Train): '+color.END+color.BOLD+str(metrics.recall_score(y_train,train_pred))+color.END)
print('\n'+color.RED+'RECALL (Test): '+color.END+color.BOLD+str(recall)+color.END)


```

```



# precision
precision = metrics.precision_score(y_test, test_pred)
print('\n'+color.RED+'PRECISION (Train) : '+color.END+color.BOLD+str(metrics.precision)
print('\n'+color.RED+'PRECISION (Test) : '+color.END+color.BOLD+str(precision)+color
(21000, 13) (9000, 13) (21000,) (9000,)

```

Confusion Matrix for Train data



Confusion Matrix for Test data



AUC (Train): 0.9358427023980465

AUC (Test): 0.9156164440716609

F1 SCORE (Train) : 0.9407600096543217

F1 SCORE (Test) : 0.9351753453772583

RECALL (Train): 0.992867655383222

RECALL (Test): 0.9903349662385807

PRECISION (Train) : 0.8938490546807318

PRECISION (Test) : 0.8858360966366651

5.5 [4.4] Word2Vec

```
In [50]: X = final['preprocessed_reviews']
        y = final['Score']

        # split the data set into train and test
        X_train, x_test, y_train, y_test = model_selection.train_test_split(X, y, test_size=0

        print(X_train.shape, x_test.shape, y_train.shape, y_test.shape)

(53089,) (22753,) (53089,) (22753,)
```

```
In [51]: # Train your own Word2Vec model using your own text corpus
```

```
# Train data
list_of_sentence=[]
for sentence in X_train:
    list_of_sentence.append(sentence.split())

# Test data
list_of_test_sentence = []
for sentence in x_test:
    list_of_test_sentence.append(sentence.split())
```

```
In [52]: # Using Google News Word2Vectors
```

```
# in this project we are using a pretrained model by google
# its 3.3G file, once you load this into your memory
# it occupies ~9Gb, so please do this step only if you have >12G of ram
# we will provide a pickle file wich contains a dict ,
```

```

# and it contains all our corpus words as keys and model[word] as values
# To use this code-snippet, download "GoogleNews-vectors-negative300.bin"
# from https://drive.google.com/file/d/0B7XkCwpI5KDYNlNUTTlSS21pQmM/edit
# it's 1.9GB in size.

# http://kavita-ganesan.com/gensim-word2vec-tutorial-starter-code/#.W17SRFAzZPY
# you can comment this whole cell
# or change these variable according to your need

is_your_ram_gt_16g=False
want_to_use_google_w2v = False
want_to_train_w2v = True

if want_to_train_w2v:
    # min_count = 5 considers only words that occurred at least 5 times
    # train data
    w2v_model_tr=Word2Vec(list_of_sentence,min_count=5,size=50, workers=4)
    # train model on test data
    w2v_model_test = Word2Vec(list_of_test_sentence,min_count=5,size=50, workers=4)
    print(w2v_model_tr.wv.most_similar('great'))
    print('='*50)
    print(w2v_model_tr.wv.most_similar('worst'))

elif want_to_use_google_w2v and is_your_ram_gt_16g:
    if os.path.isfile('GoogleNews-vectors-negative300.bin'):
        w2v_model=KeyedVectors.load_word2vec_format('GoogleNews-vectors-negative300.b
        print(w2v_model.wv.most_similar('great'))
        print(w2v_model.wv.most_similar('worst'))
    else:
        print("you don't have google's word2vec file, keep want_to_train_w2v = True, t

[('fantastic', 0.8369816541671753), ('awesome', 0.83486008644104), ('good', 0.8010945916175842)
=====
[('greatest', 0.7605360746383667), ('best', 0.6946088075637817), ('tastiest', 0.69453716278076

In [53]: # train data operation
w2v_train_words = list(w2v_model_tr.wv.vocab)
print("number of words that occurred minimum 5 times ",len(w2v_train_words))
print("sample words ", w2v_train_words[0:50])

number of words that occurred minimum 5 times 13914
sample words ['product', 'china', 'known', 'would', 'not', 'ordered', 'share', 'concerns', 'f

In [54]: ## test data operation
w2v_test_words = list(w2v_model_test.wv.vocab)

```

```

print("number of words that occurred minimum 5 times ",len(w2v_test_words))
print("sample words ", w2v_test_words[0:50])

```

number of words that occurred minimum 5 times 9272

sample words ['recently', 'recieved', 'samples', 'energy', 'husband', 'son', 'used', 'next',

5.6 [4.4.1] Converting text into vectors using Avg W2V, TFIDF-W2V

[4.4.1.1] Avg W2v

In [55]: # average Word2Vec

```

# train data operation
exists = os.path.isfile(avg_w2v_trained_model_100000)
exists = False
if exists:
    print("yes exist")
    final_w2v_train = load(avg_w2v_trained_model_100000)
else:
    print("not exist")
    # compute average word2vec for each review.
    final_w2v_train = []; # the avg-w2v for each sentence/review is stored in this list
    for sent in tqdm(list_of_sentence): # for each review/sentence
        sent_vec = np.zeros(50) # as word vectors are of zero length 50, you might need to initialize
        cnt_words = 0; # num of words with a valid vector in the sentence/review
        for word in sent: # for each word in a review/sentence
            if word in w2v_train_words:
                vec = w2v_model_tr.wv[word]
                sent_vec += vec
                cnt_words += 1
        if cnt_words != 0:
            sent_vec /= cnt_words
        final_w2v_train.append(sent_vec)
    print(len(final_w2v_train))
    print(len(final_w2v_train[0]))
# dump(final_w2v_train, avg_w2v_trained_model_100000)

# test data operation
exists = os.path.isfile(avg_w2v_test_model_100000)
exists = False
if exists:
    print("yes exist")
    final_w2v_test = load(avg_w2v_test_model_100000)
else:
    print("not exist")
    final_w2v_test = []; # the avg-w2v for each sentence/review is stored in this list
    for sent in tqdm(list_of_test_sentence): # for each review/sentence

```



```

sent_vec = np.zeros(50) # as word vectors are of zero length 50, you might ne
cnt_words =0; # num of words with a valid vector in the sentence/review
for word in sent: # for each word in a review/sentence
    if word in w2v_test_words:
        vec = w2v_model_test.wv[word]
        sent_vec += vec
        cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    final_w2v_test.append(sent_vec)
print(len(final_w2v_test))
print(len(final_w2v_test[0]))
#     dump(final_w2v_test, avg_w2v_test_model_100000)

```

```
0%|          | 0/53089 [00:00<?, ?it/s]
```

not exist

```
100%|| 53089/53089 [01:42<00:00, 517.62it/s]
 4%|          | 926/22753 [00:01<00:25, 870.92it/s]
```

53089

50

not exist

```
100%|| 22753/22753 [00:30<00:00, 754.78it/s]
```

22753

50

5.6.1 Hyper param Tuning using GridSearch

finding 'max depth' which have maximum AUC Score

```
In [56]: w2v_train_path = '/home/pranay/ML Hyperparam Tune/GBDT/w2v_train_hyperparam_tuned'
exists = os.path.isfile(w2v_train_path)
```

```

if exists:
    print("yes exists")
    w2v_train = load(w2v_train_path)
else:
    print("not exists")
    final_w2v_train = np.asarray(final_w2v_train)

```

```

w2v_train = finding_best_hyperparam(final_w2v_train,y_train)
dump(w2v_train,w2v_train_path )

# view the complete results (list of named tuples)
print("=====Training=====")
print (w2v_train.best_score_)
print (w2v_train.best_params_)
print (w2v_train.best_estimator_)

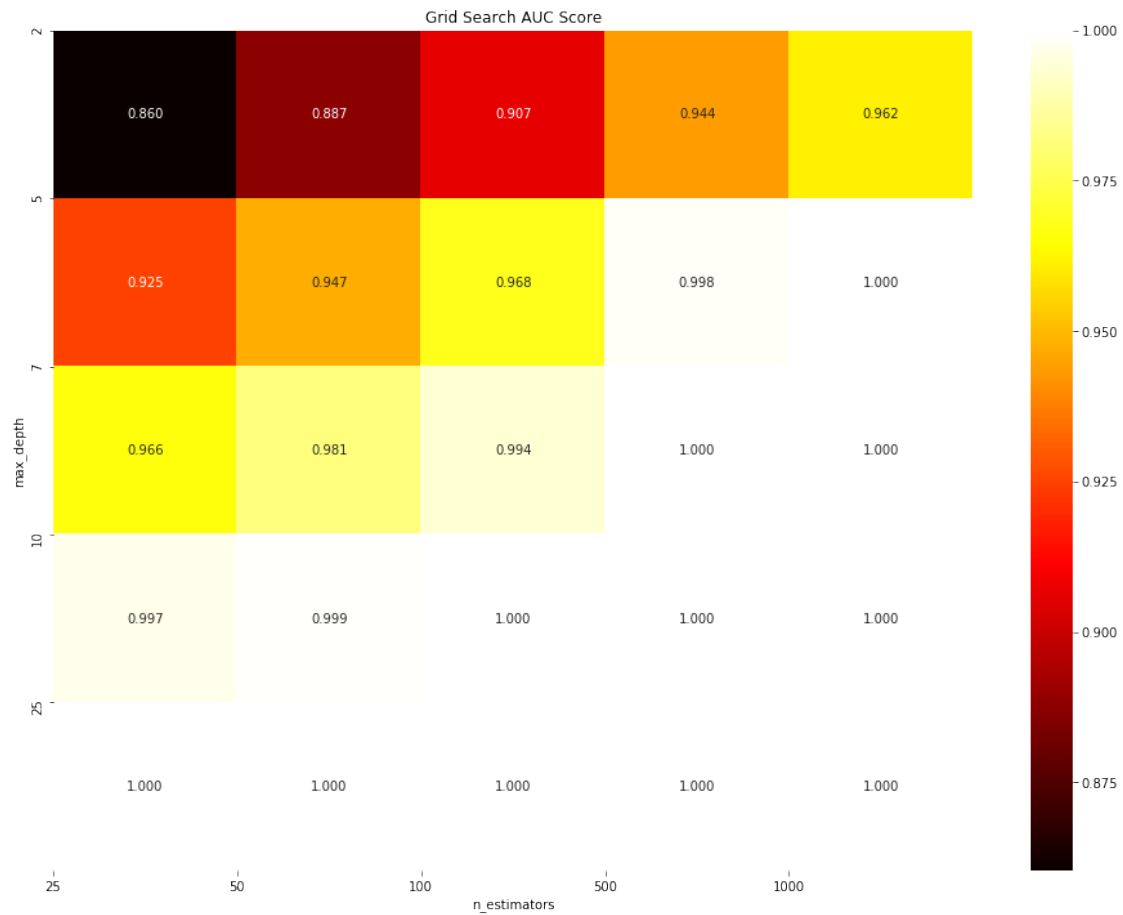
best_depth_size = w2v_train.best_params_.get("max_depth", "")
best_estimators = w2v_train.best_params_.get("n_estimators", "")

yes exists
=====Training=====
0.9010064822913382
{'max_depth': 25, 'n_estimators': 1000}
XGBClassifier(base_score=0.5, booster='gbtree', class_weight='balanced',
              colsample_bylevel=1, colsample_bytree=1, gamma=0, learning_rate=0.1,
              max_delta_step=0, max_depth=25, min_child_weight=1, missing=nan,
              n_estimators=1000, n_jobs=1, nthread=None,
              objective='binary:logistic', random_state=1, reg_alpha=0,
              reg_lambda=1, scale_pos_weight=1, seed=None, silent=True,
              subsample=1)

In [57]: print('\n'+color.BOLD +'AUC Train data' +color.END)
         plotHeatMap(w2v_train,'trained')


```

AUC Train data



```
In [58]: print('\n'+color.BOLD + 'AUC Validation data'+color.END)
         plotHeatMap(w2v_train, 'test')
```

AUC Validation data



```
In [77]: final_w2v_test = np.asarray(final_w2v_test)
         final_w2v_train = np.asarray(final_w2v_train)
```

```
optimal_model = XGBClassifier(base_score=0.5, booster='gbtree', class_weight='balanced',
                             colsample_bylevel=1, colsample_bytree=1, gamma=0, learning_rate=0.1,
                             max_delta_step=0, max_depth=3, min_child_weight=1, missing=None,
                             n_estimators=100, n_jobs=-1, nthread=None,
                             objective='binary:logistic', random_state=1, reg_alpha=0,
                             reg_lambda=1, scale_pos_weight=1, seed=None, silent=True,
                             subsample=1)
```

```
# fitting the model
```

```
optimal_model.fit(final_w2v_train, y_train)
```

```
# predict the response
```

```
test_pred = optimal_model.predict(final_w2v_test)
```

```
train_pred = optimal_model.predict(final_w2v_train)
```

```

print('\n'+color.RED+'Max Depth : '+color.END+color.BOLD+str(3)+color.END)
print('\n'+color.RED+'Estimators : '+color.END+color.BOLD+str(100)+color.END)

# plot confusion matrix

print('\n'+color.BOLD+'Confusion Matrix for Train data'+color.END)
plotConfusionMatrix(y_train,train_pred)

print('\n'+color.BOLD+'Confusion Matrix for Test data'+color.END)
plotConfusionMatrix(y_test,test_pred)

# plot AUC
train_auc,test_auc = plotAUC_ROC(optimal_model,final_w2v_train, y_train,final_w2v_test)
print('\n'+color.RED+'AUC (Train): '+color.END+color.BOLD+str(train_auc)+color.END)
print('\n'+color.RED+'AUC (Test): '+color.END+color.BOLD+str(test_auc)+color.END)

# f1 score
score = f1_score(y_test,test_pred)
print('\n'+color.RED+'F1 SCORE (Train) : '+color.END+color.BOLD+str(f1_score(y_train,train_pred))+color.END)
print('\n'+color.RED+'F1 SCORE (Test) : '+color.END+color.BOLD+str(score)+color.END)

# recall
recall = metrics.recall_score(y_test, test_pred)
print('\n'+color.RED+'RECALL (Train): '+color.END+color.BOLD+str(metrics.recall_score(y_train,train_pred))+color.END)
print('\n'+color.RED+'RECALL (Test): '+color.END+color.BOLD+str(recall)+color.END)


# precision
precision = metrics.precision_score(y_test, test_pred)
print('\n'+color.RED+'PRECISION (Train) : '+color.END+color.BOLD+str(metrics.precision_score(y_train,train_pred))+color.END)
print('\n'+color.RED+'PRECISION (Test) : '+color.END+color.BOLD+str(precision)+color.END)

```


Max Depth : 3


Estimators : 100

Confusion Matrix for Train data



Confusion Matrix for Test data





AUC (Train): 0.9117189542570628

AUC (Test): 0.8294383503755657

F1 SCORE (Train) : 0.9354866494401377

F1 SCORE (Test) : 0.8972166153192394

RECALL (Train): 0.976928266246908

RECALL (Test): 0.8895149823582074

PRECISION (Train) : 0.8974178888659368

PRECISION (Test) : 0.9050527782242941

5.7 [4.4.1.2] TFIDF weighted W2v

```
In [61]: X = final['preprocessed_reviews']
         y = final['Score']
```

```

# split the data set into train and test
X_train, x_test, y_train, y_test = model_selection.train_test_split(X, y, test_size=0.2)

print(X_train.shape, x_test.shape, y_train.shape, y_test.shape)

(53089,) (22753,) (53089,) (22753,)

In [62]: # S = ["abc def pqr", "def def def abc", "pqr pqr def"]
model = TfidfVectorizer()
tf_idf_matrix = model.fit_transform(preprocessed_reviews)
# we are converting a dictionary with word as a key, and the idf as a value
dictionary = dict(zip(model.get_feature_names(), list(model.idf_)))

In [63]: # we are converting a dictionary with word as a key, and the idf as a value
dictionary = dict(zip(tf_idf_vect.get_feature_names(), list(tf_idf_vect.idf_)))

# TF-IDF weighted Word2Vec

# Train data operation

# store model to hard disk if exist then load model directly from memory
exists = os.path.isfile(w2v_tf_idf_trained_model_100000)
exists = False
if exists:
    print("yes exist")
    final_tfidf_w2v_tr = load(w2v_tf_idf_trained_model_100000)
else:
    print("not exist")
    tfidf_feat = tf_idf_vect.get_feature_names() # tfidf words/col-names
    # final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

    final_tfidf_w2v_tr = []; # the tfidf-w2v for each sentence/review is stored in the list
    row=0;
    for sent in tqdm(list_of_sentence): # for each review/sentence
        sent_vec = np.zeros(50) # as word vectors are of zero length
        weight_sum = 0; # num of words with a valid vector in the sentence/review
        for word in sent: # for each word in a review/sentence
            if word in w2v_train_words and word in tfidf_feat:
                vec = w2v_model_tr.wv[word]
                # tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
                # to reduce the computation we are
                # dictionary[word] = idf value of word in whole corpus
                # sent.count(word) = tf value of word in this review
                tf_idf = dictionary[word]*(sent.count(word)/len(sent))
                sent_vec += (vec * tf_idf)
            weight_sum += tf_idf

```



```

        if weight_sum != 0:
            sent_vec /= weight_sum
            final_tfidf_w2v_tr.append(sent_vec)
            row += 1
#         dump(final_tfidf_w2v_tr, w2v_tf_idf_trained_model_100000)

# Test data operation =====

# store model to hard disk if exist then load model directly from memory
exists = os.path.isfile(w2v_tf_idf_test_model_100000)
exists = False
if exists:
    print("yes exist")
    final_tfidf_w2v_test = load(w2v_tf_idf_test_model_100000)

else:
    print("not exist")
    # final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = t
    final_tfidf_w2v_test = []; # the tfidf-w2v for each sentence/review is stored in
    row=0;
    for sent in tqdm(list_of_test_sentence): # for each review/sentence
        sent_vec = np.zeros(50) # as word vectors are of zero length
        weight_sum = 0; # num of words with a valid vector in the sentence/review
        for word in sent: # for each word in a review/sentence
            if word in w2v_test_words and word in tfidf_feat:
                vec = w2v_model_test.wv[word]
            #         tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
            # to reduce the computation we are
            # dictionary[word] = idf value of word in whole corpus
            # sent.count(word) = tf value of word in this review
            tf_idf = dictionary[word]*(sent.count(word)/len(sent))
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
        if weight_sum != 0:
            sent_vec /= weight_sum
            final_tfidf_w2v_test.append(sent_vec)
            row += 1
#         dump(final_tfidf_w2v_test, w2v_tf_idf_test_model_100000)

```

```
0%|          | 7/53089 [00:00<28:05, 31.50it/s]
```

```
not exist
```

```
100%|| 53089/53089 [43:17<00:00, 21.13it/s]
0%|          | 32/22753 [00:01<19:51, 19.07it/s]
```

not exist

100%|| 22753/22753 [16:18<00:00, 23.25it/s]

5.7.1 Hyper param Tuning using GridSearch

finding 'max depth' which have maximum AUC Score

```
In [64]: tfidf_w2v_train_path = '/home/pranay/ML Hyperparam Tune/GBDT/tfidf_w2v_train_hyperparam.pkl'
exists = os.path.isfile(tfidf_w2v_train_path)
```

```
if exists:
    print("yes exists")
    w2v_tfidf_train = load(tfidf_w2v_train_path)
else:
    print("not exists")
    final_tfidf_w2v_tr = np.asarray(final_tfidf_w2v_tr)
    w2v_tfidf_train = finding_best_hyperparam(final_tfidf_w2v_tr,y_train)
    dump(w2v_tfidf_train,tfidf_w2v_train_path )
```

```
# view the complete results (list of named tuples)
```

```
print("====Training====")
print(w2v_tfidf_train.best_score_)
print(w2v_tfidf_train.best_params_)
print(w2v_tfidf_train.best_estimator_)
```

```
best_depth_size = w2v_tfidf_train.best_params_.get("max_depth", "")
best_estimators = w2v_tfidf_train.best_params_.get("n_estimators", "")
```

yes exists

====Training====


0.8782853837864152

{'max_depth': 10, 'n_estimators': 500}

XGBClassifier(base_score=0.5, booster='gbtree', class_weight='balanced',
colsample_bylevel=1, colsample_bytree=1, gamma=0, learning_rate=0.1,
max_delta_step=0, max_depth=10, min_child_weight=1, missing=nan,
n_estimators=500, n_jobs=1, nthread=None,
objective='binary:logistic', random_state=1, reg_alpha=0,
reg_lambda=1, scale_pos_weight=1, seed=None, silent=True,
subsample=1)


```
In [65]: print('\n'+color.BOLD + 'AUC Train data'+color.END)
plotHeatMap(w2v_tfidf_train,'trained')
```

AUC Train data



```
In [66]: print('\n'+color.BOLD + 'AUC Validation data'+color.END)
          plotHeatMap(w2v_tfidf_train, 'validation')
```

AUC Validation data



5.8 GBDT on TFIDF - W2V

```
In [76]: final_tfidf_w2v_test = np.asarray(final_tfidf_w2v_test)
        final_tfidf_w2v_tr = np.asarray(final_tfidf_w2v_tr)
```

```
optimal_model = XGBClassifier(base_score=0.5, booster='gbtree', class_weight='balanced',
                             colsample_bylevel=1, colsample_bytree=1, gamma=0, learning_rate=0.1,
                             max_delta_step=0, max_depth=2, min_child_weight=1, missing=None,
                             n_estimators=80, n_jobs=-1, nthread=None,
                             objective='binary:logistic', random_state=1, reg_alpha=0,
                             reg_lambda=1, scale_pos_weight=1, seed=None, silent=True,
                             subsample=1)
```

```
# fitting the model
```

```
optimal_model.fit(final_tfidf_w2v_tr, y_train)
```

```
# predict the response
```

```
test_pred = optimal_model.predict(final_tfidf_w2v_test)
```

```
train_pred = optimal_model.predict(final_tfidf_w2v_tr)
```

```

print('\n'+color.RED+'Max Depth : '+color.END+color.BOLD+str(2)+color.END)
print('\n'+color.RED+'Estimators: '+color.END+color.BOLD+str(80)+color.END)

# plot confusion matrix

print('\n'+color.BOLD+'Confusion Matrix for Train data'+color.END)
plotConfusionMatrix(y_train,train_pred)

print('\n'+color.BOLD+'Confusion Matrix for Test data'+color.END)
plotConfusionMatrix(y_test,test_pred)

# plot AUC
train_auc,test_auc = plotAUC_ROC(optimal_model,final_tfidf_w2v_tr, y_train,final_tfidf)
print('\n'+color.RED+'AUC (Train): '+color.END+color.BOLD+str(train_auc)+color.END)
print('\n'+color.RED+'AUC (Test): '+color.END+color.BOLD+str(test_auc)+color.END)

# f1 score
score = f1_score(y_test,test_pred)
print('\n'+color.RED+'F1 SCORE (Train) : '+color.END+color.BOLD+str(f1_score(y_train,train_pred))+color.END)
print('\n'+color.RED+'F1 SCORE (Test) : '+color.END+color.BOLD+str(score)+color.END)

# recall
recall = metrics.recall_score(y_test, test_pred)
print('\n'+color.RED+'RECALL (Train): '+color.END+color.BOLD+str(metrics.recall_score(y_train,train_pred))+color.END)
print('\n'+color.RED+'RECALL (Test): '+color.END+color.BOLD+str(recall)+color.END)


# precision
precision = metrics.precision_score(y_test, test_pred)
print('\n'+color.RED+'PRECISION (Train) : '+color.END+color.BOLD+str(metrics.precision_score(y_train,train_pred))+color.END)
print('\n'+color.RED+'PRECISION (Test) : '+color.END+color.BOLD+str(precision)+color.END)

```


Max Depth : 2

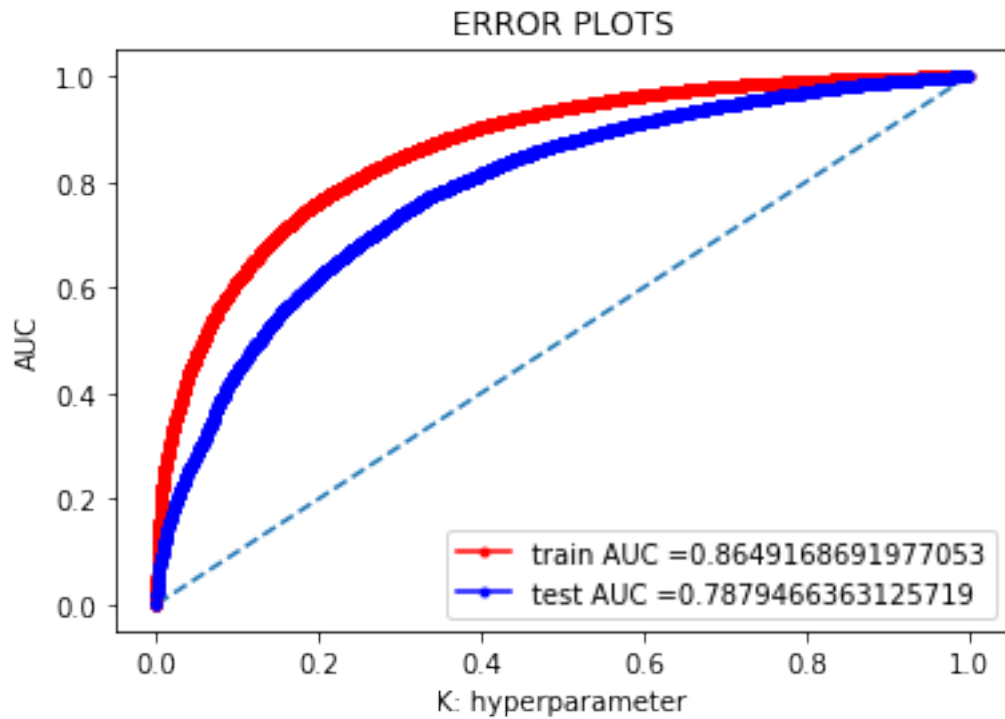
Estimators: 80

Confusion Matrix for Train data



Confusion Matrix for Test data





AUC (Train): 0.8649168691977053

AUC (Test): 0.7879466363125719

F1 SCORE (Train) : 0.9244121293630614

F1 SCORE (Test) : 0.9124840388959827

RECALL (Train): 0.989206206431302

RECALL (Test): 0.9784612143872768

PRECISION (Train) : 0.8675844114862733

PRECISION (Test) : 0.8548424200598114

6 [6] Conclusions

```
In [84]: import pandas as pd
         from prettytable import PrettyTable

         print(color.BOLD+'\t\t\t\t\t GBDT \t'+color.END)
         print('\n')

         print(color.BOLD+'For BOW and TFIDF, We have considered 85k points'+color.END)
         print(color.BOLD+'For BOW- Additional Feature and TFIDF- Additional Feature, We have c

         x = PrettyTable()
         x.field_names = ['Metric', 'BOW', 'BOW-Additional Feature', 'TFIDF', 'TFIDF- Additional

         x.add_row(["Max Depth ", 4,4,3,3,3,2])
         x.add_row(["no Estimator ", 100,100,100,100,100,80])

         x.add_row(["AUC Train ", 0.91487,0.99454,0.90289,0.93584,0.91171,0.86491])
         x.add_row(["AUC Test ", 0.90164,0.96708,0.88776,0.91561,0.82943,0.78794])

         x.add_row(["F1 SCORE Train ", 0.93603,0.98221,0.93143,0.94076,0.93548,0.92441])
         x.add_row(["F1 SCORE Test ", 0.93118,0.96068,0.92721,0.93517,0.89721,0.91248])

         x.add_row(["RECALL Train ",0.99554,0.99575,0.99707,0.99286,0.97692,0.98920])
         x.add_row(["RECALL Test ", 0.99294,0.98345,0.99483,0.99033,0.88951,0.97846])
         93449
         x.add_row(["PRECISION Train ", 0.88323,0.96904,0.87390,0.89384,0.89741,0.86751])
         x.add_row(["PRECISION Test ",0.87665,0.93894,0.86819,0.88583,0.90502,0.85484])

         print('\n')
         print(x)
```

GBDT

For BOW and TFIDF, We have considered 85k points

For BOW- Additional Feature and TFIDF- Additional Feature, We have considered 30k points

Metric	BOW	BOW-Additional Feature	TFIDF	TFIDF- Additional Features
Max Depth	4	4	3	3
no Estimator	100	100	100	100
AUC Train	0.91487	0.99454	0.90289	0.93584

AUC Test	0.90164	0.96708	0.88776	0.91561	
F1 SCORE Train	0.93603	0.98221	0.93143	0.94076	
F1 SCORE Test	0.93118	0.96068	0.92721	0.93517	
RECALL Train	0.99554	0.99575	0.99707	0.99286	
RECALL Test	0.99294	0.98345	0.99483	0.99033	
PRECISION Train	0.88323	0.96904	0.8739	0.89384	
PRECISION Test	0.87665	0.93894	0.86819	0.88583	
+-----+-----+-----+-----+-----+					

In []: