

# Implement SGD

June 18, 2019

```
In [ ]: ## reference link : https://spin.atomicobject.com/2014/06/24/gradient-descent-linear-r
```

```
In [23]: # Imported necessary libraries
from sklearn.datasets import load_boston
from sklearn.model_selection import train_test_split
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics import mean_squared_error
from sklearn import preprocessing
from sklearn.linear_model import SGDRegressor

# Data loaded
boston = load_boston()

# Data shape
boston.data.shape
```

```
Out[23]: (506, 13)
```

```
In [24]: # Feature name
boston.feature_names
```

```
Out[24]: array(['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS', 'RAD',
               'TAX', 'PTRATIO', 'B', 'LSTAT'], dtype='<U7')
```

```
In [25]: # This is y value i.e. target
boston.target.shape
```

```
Out[25]: (506,)
```

```
In [26]: # Convert it into pandas dataframe
data = pd.DataFrame(boston.data, columns = boston.feature_names)
data.head()
```

```
Out[26]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	\
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	

2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0

	PTRATIO	B	LSTAT
0	15.3	396.90	4.98
1	17.8	396.90	9.14
2	17.8	392.83	4.03
3	18.7	394.63	2.94
4	18.7	396.90	5.33

```
In [27]: # Statistical summary
data.describe()
```

```
Out[27]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	\
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	
mean	3.613524	11.363636	11.136779	0.069170	0.554695	6.284634	
std	8.601545	23.322453	6.860353	0.253994	0.115878	0.702617	
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500	
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500	
75%	3.677083	12.500000	18.100000	0.000000	0.624000	6.623500	
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	

	AGE	DIS	RAD	TAX	PTRATIO	B	\
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	
mean	68.574901	3.795043	9.549407	408.237154	18.455534	356.674032	
std	28.148861	2.105710	8.707259	168.537116	2.164946	91.294864	
min	2.900000	1.129600	1.000000	187.000000	12.600000	0.320000	
25%	45.025000	2.100175	4.000000	279.000000	17.400000	375.377500	
50%	77.500000	3.207450	5.000000	330.000000	19.050000	391.440000	
75%	94.075000	5.188425	24.000000	666.000000	20.200000	396.225000	
max	100.000000	12.126500	24.000000	711.000000	22.000000	396.900000	

	LSTAT
count	506.000000
mean	12.653063
std	7.141062
min	1.730000
25%	6.950000
50%	11.360000
75%	16.955000
max	37.970000

```
In [28]: #standardize for fast convergence to minima
data = (data - data.mean())/data.std()
data.head()
```

```
Out[28]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	\
0	-0.419367	0.284548	-1.286636	-0.272329	-0.144075	0.413263	-0.119895	

```

1 -0.416927 -0.487240 -0.592794 -0.272329 -0.739530 0.194082 0.366803
2 -0.416929 -0.487240 -0.592794 -0.272329 -0.739530 1.281446 -0.265549
3 -0.416338 -0.487240 -1.305586 -0.272329 -0.834458 1.015298 -0.809088
4 -0.412074 -0.487240 -1.305586 -0.272329 -0.834458 1.227362 -0.510674

```

```

      DIS      RAD      TAX  PTRATIO      B      LSTAT
0  0.140075 -0.981871 -0.665949 -1.457558 0.440616 -1.074499
1  0.556609 -0.867024 -0.986353 -0.302794 0.440616 -0.491953
2  0.556609 -0.867024 -0.986353 -0.302794 0.396035 -1.207532
3  1.076671 -0.752178 -1.105022 0.112920 0.415751 -1.360171
4  1.076671 -0.752178 -1.105022 0.112920 0.440616 -1.025487

```

```

In [29]: # MEDV(median value is usually target), change it to price
data["PRICE"] = boston.target
data.head()

```

```

Out[29]:      CRIM      ZN      INDUS      CHAS      NOX      RM      AGE \
0 -0.419367 0.284548 -1.286636 -0.272329 -0.144075 0.413263 -0.119895
1 -0.416927 -0.487240 -0.592794 -0.272329 -0.739530 0.194082 0.366803
2 -0.416929 -0.487240 -0.592794 -0.272329 -0.739530 1.281446 -0.265549
3 -0.416338 -0.487240 -1.305586 -0.272329 -0.834458 1.015298 -0.809088
4 -0.412074 -0.487240 -1.305586 -0.272329 -0.834458 1.227362 -0.510674

```

```

      DIS      RAD      TAX  PTRATIO      B      LSTAT  PRICE
0  0.140075 -0.981871 -0.665949 -1.457558 0.440616 -1.074499 24.0
1  0.556609 -0.867024 -0.986353 -0.302794 0.440616 -0.491953 21.6
2  0.556609 -0.867024 -0.986353 -0.302794 0.396035 -1.207532 34.7
3  1.076671 -0.752178 -1.105022 0.112920 0.415751 -1.360171 33.4
4  1.076671 -0.752178 -1.105022 0.112920 0.440616 -1.025487 36.2

```

```

In [30]: # Target and features
Y = data["PRICE"]
X = data.drop("PRICE", axis = 1)

```

```

In [31]: from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(X, Y, test_size = 0.3)
print(x_train.shape, x_test.shape, y_train.shape, y_test.shape)

x_train["PRICE"] = y_train

```

```

(354, 13) (152, 13) (354,) (152,)

```

/home/pranay/anaconda3/lib/python3.7/site-packages/ipykernel\_launcher.py:5: SettingWithCopyWarning: A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html>

## 0.1 Custome SGD Implementation

```
In [32]: def gradient_decent(w0, b0, points, x_test, y_test, learning_rate):
    iterations = 1000
    gradient_m = 0
    gradient_b = 0
    cost_train = []
    cost_test = []
    for j in range(1, iterations):
```

```
        # Random Train
```

```
        train_sample = points.sample(160)
        y = np.asmatrix(train_sample["PRICE"])
        x = np.asmatrix(train_sample.drop("PRICE", axis = 1))
```

```
        for i in range(len(x)):
            dot_prod = np.dot(x[i] , w0) + b0
            subtract_term = y[:,i] - dot_prod
            gradient_m += np.dot(-2*x[i].T , subtract_term)
            gradient_b += -2*(subtract_term)
```

```
        # print(subtract_term.shape)
```

```
        w1 = w0 - learning_rate * gradient_m
        b1 = b0 - learning_rate * gradient_b
```

```
        if (w0==w1).all():
            break
        else:
            w0 = w1
            b0 = b1
            learning_rate = learning_rate/2
```

```
    return w0, b0
```

```
def mse_metric(b, m, features, target):
    totalError = 0
    for i in range(0, len(features)):
        x = features
        y = target
        totalError += (y[:,i] - (np.dot(x[i] , m) + b)) ** 2
    return totalError / len(x)
```

```
In [33]: learning_rate = 0.001
    w0_random = np.random.rand(13)
    w0 = np.asmatrix(w0_random).T
    b0 = np.random.rand()
```

```
    optimal_w, optimal_b = gradient_decent(w0, b0, x_train, x_test, y_test, learning_rate)
    print("Coefficient: {} \n y_intercept: {}".format(optimal_w, optimal_b))
```


```
Coefficient: [[-0.68292829]
 [ 1.13355708]
 [-0.13956617]
 [ 3.11930798]
 [-0.24116852]
 [ 4.15839783]
 [ 0.28884544]
 [-0.66286099]
 [ 0.16888711]
 [-0.34984621]
 [ 0.06535574]
 [ 1.62262885]
 [-3.5635532 ]]
y_intercept: [[20.65043429]]
```

```
In [34]: # Implemented SGD
        # The mean squared error
        mse_error = mse_metric(optimal_b, optimal_w, np.asmatrix(x_test), np.asmatrix(y_test))
        print("Custom SGD Mean squared error: %.2f" % mse_error)
```

Custom SGD Mean squared error: 34.08

```
In [35]: y_pred= [(np.dot(np.asmatrix(x_test), optimal_w) + optimal_b)]

        # manual_error=plot_(y_pred)
        plt.scatter(y_test, y_pred)
        plt.xlabel("Prices:  $Y_i$ ")
        plt.ylabel("Predicted prices:  $\hat{Y}_i$ ")
        plt.title("Prices vs Predicted prices:  $Y_i$  vs  $\hat{Y}_i$ ")
        plt.show()
```



## 0.2 SKLearn SGD Implementation

```
In [36]: X = load_boston().data
        Y = load_boston().target
```

```
df=pd.DataFrame(X)
df.head()
```

```
# standardise data
scaler = preprocessing.StandardScaler()
X = scaler.fit_transform(X)
df=pd.DataFrame(X)
```

```
df['price']=Y
df.head()
```

```
Out [36]:
```

	0	1	2	3	4	5	6 \
0	-0.419782	0.284830	-1.287909	-0.272599	-0.144217	0.413672	-0.120013
1	-0.417339	-0.487722	-0.593381	-0.272599	-0.740262	0.194274	0.367166
2	-0.417342	-0.487722	-0.593381	-0.272599	-0.740262	1.282714	-0.265812
3	-0.416750	-0.487722	-1.306878	-0.272599	-0.835284	1.016303	-0.809889
4	-0.412482	-0.487722	-1.306878	-0.272599	-0.835284	1.228577	-0.511180

	7	8	9	10	11	12	price
0	0.140214	-0.982843	-0.666608	-1.459000	0.441052	-1.075562	24.0
1	0.557160	-0.867883	-0.987329	-0.303094	0.441052	-0.492439	21.6
2	0.557160	-0.867883	-0.987329	-0.303094	0.396427	-1.208727	34.7
3	1.077737	-0.752922	-1.106115	0.113032	0.416163	-1.361517	33.4
4	1.077737	-0.752922	-1.106115	0.113032	0.441052	-1.026501	36.2

```
In [37]: # standardise data
scaler = preprocessing.StandardScaler()
X = scaler.fit_transform(X)
df=pd.DataFrame(X)

df['price']=Y
df.head()
```

```
Out [37]:
```

	0	1	2	3	4	5	6	\
0	-0.419782	0.284830	-1.287909	-0.272599	-0.144217	0.413672	-0.120013	
1	-0.417339	-0.487722	-0.593381	-0.272599	-0.740262	0.194274	0.367166	
2	-0.417342	-0.487722	-0.593381	-0.272599	-0.740262	1.282714	-0.265812	
3	-0.416750	-0.487722	-1.306878	-0.272599	-0.835284	1.016303	-0.809889	
4	-0.412482	-0.487722	-1.306878	-0.272599	-0.835284	1.228577	-0.511180	

	7	8	9	10	11	12	price
0	0.140214	-0.982843	-0.666608	-1.459000	0.441052	-1.075562	24.0
1	0.557160	-0.867883	-0.987329	-0.303094	0.441052	-0.492439	21.6
2	0.557160	-0.867883	-0.987329	-0.303094	0.396427	-1.208727	34.7
3	1.077737	-0.752922	-1.106115	0.113032	0.416163	-1.361517	33.4
4	1.077737	-0.752922	-1.106115	0.113032	0.441052	-1.026501	36.2

```
In [38]: # Split data into train and test
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = 0.33, random_st
print(X_train.shape)
print(X_test.shape)
print(Y_train.shape)
print(Y_test.shape)
```

```
(339, 13)
(167, 13)
(339,)
(167,)
```

```
In [39]: from sklearn.linear_model import SGDRegressor
from sklearn.metrics import mean_squared_error, r2_score
clf = SGDRegressor()
clf.fit(X_train, Y_train)
Y_pred = clf.predict(X_test)

print("Coefficients: \n", clf.coef_)
print("Y_intercept", clf.intercept_)
```

Coefficients:

```
[-0.9266483  0.19881195 -0.6338864  0.27567329 -0.45419415  3.06297705  
-0.40894034 -1.82317767  1.04266262 -0.41628519 -1.98692311  0.97787009  
-3.06916569]
```


Y\_intercept [21.75067408]

```
/home/pranay/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/stochastic_gradient.py  
FutureWarning)
```

```
In [40]: # Sklearn SGD  
# The mean squared error  
print("SKLearn Mean squared error: %.2f" % mean_squared_error(Y_test, Y_pred))
```

SKLearn Mean squared error: 30.59

```
In [41]: plt.scatter(Y_test, Y_pred)  
plt.xlabel("Prices: $Y_i$")  
plt.ylabel("Predicted prices: $\hat{Y}_i$")  
plt.title("Prices vs Predicted prices: $Y_i$ vs $\hat{Y}_i$")  
plt.show()
```





```
In [45]: import pandas as pd
         from prettytable import PrettyTable

         x = PrettyTable()
         x.field_names = ['Metric', 'Sklearn', 'Custokm SGD']
         x.add_row(["MSE vales ", 30.59, 34.08])

         print('\n')
         print(x)
```

```
+-----+-----+-----+
| Metric | Sklearn | Custokm SGD |
+-----+-----+-----+
| MSE vales | 30.59 | 34.08 |
+-----+-----+-----+
```

## 0.2.1 final weight from Custom SGD

```
In [43]: print("Coefficient: {} \n y_intercept: {}".format(optimal_w, optimal_b))
```

```
Coefficient: [[-0.68292829]
 [ 1.13355708]
 [-0.13956617]
 [ 3.11930798]
 [-0.24116852]
 [ 4.15839783]
 [ 0.28884544]
 [-0.66286099]
 [ 0.16888711]
 [-0.34984621]
 [ 0.06535574]
 [ 1.62262885]
 [-3.5635532 ]]
y_intercept: [[20.65043429]]
```

## 0.2.2 final weight from SGD Sklearn

```
In [44]: print("Coefficients: \n", clf.coef_)
         print("Y_intercept", clf.intercept_)
```

```
Coefficients:
[-0.9266483  0.19881195 -0.6338864  0.27567329 -0.45419415  3.06297705
-0.40894034 -1.82317767  1.04266262 -0.41628519 -1.98692311  0.97787009
-3.06916569]
```

Y\_intercept [21.75067408]