

MNIST using Keras

June 24, 2019

0.1 Keras -- MLPs on MNIST

```
In [0]: # if you keras is not using tensorflow as backend set "KERAS_BACKEND=tensorflow" use t
from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal

from keras.models import Sequential
from keras.layers import Dense, Activation
from keras.layers.normalization import BatchNormalization
from keras.layers import Dropout

In [0]: %matplotlib notebook
%matplotlib inline

import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()
    plt.show()

In [3]: # the data, shuffled and split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()

Downloading data from https://s3.amazonaws.com/img-datasets/mnist.npz
11493376/11490434 [=====] - 1s 0us/step

In [4]: print("Number of training examples :", X_train.shape[0], "and each image is of shape (")
print("Number of training examples :", X_test.shape[0], "and each image is of shape (")
```

Number of training examples : 60000 and each image is of shape (28, 28)
 Number of training examples : 10000 and each image is of shape (28, 28)

```
In [0]: # if you observe the input shape its 2 dimensional vector
        # for each image we have a (28*28) vector
        # we will convert the (28*28) vector into single dimensional vector of 1 * 784
```

```
X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

```
In [6]: # after converting the input images from 3d to 2d vectors
```

```
print("Number of training examples :", X_train.shape[0], "and each image is of shape (%)")
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%)")
```

Number of training examples : 60000 and each image is of shape (784)
 Number of training examples : 10000 and each image is of shape (784)

```
In [7]: # An example data point
        print(X_train[0])
```

```
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  3  18  18  18 126 136 175  26 166 255
247 127  0  0  0  0  0  0  0  0  0  0  0  0  0  30  36  94 154
170 253 253 253 253 225 172 253 242 195  64  0  0  0  0  0  0
  0  0  0  0  0  49 238 253 253 253 253 253 253 253 251  93  82
 82  56  39  0  0  0  0  0  0  0  0  0  0  0  0  18 219 253
253 253 253 253 198 182 247 241  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  80 156 107 253 253 205  11  0  43 154
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0 14  1 154 253  90  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0 139 253 190  2  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0 11 190 253  70  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  35 241
225 160 108  1  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  81 240 253 253 119  25  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  45 186 253 253 150 27  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  16  93 252 253 187
```

```

0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 249 253 249 64 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 46 130 183 253
253 207 2 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 39 148 229 253 253 253 250 182 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 24 114 221 253 253 253
253 201 78 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 23 66 213 253 253 253 253 198 81 2 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 18 171 219 253 253 253 253 195
80 9 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
55 172 226 253 253 253 253 244 133 11 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 136 253 253 253 212 135 132 16
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0

```

```

In [0]: # if we observe the above matrix each cell is having a value between 0-255
        # before we move to apply machine learning algorithms lets try to normalize the data
        #  $X \Rightarrow (X - X_{min}) / (X_{max} - X_{min}) = X / 255$ 

```

```

X_train = X_train/255
X_test = X_test/255

```

```

In [9]: # here we are having a class number for each image
        print("Class label of first image :", y_train[0])

        # lets convert this into a 10 dimensional vector
        # ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
        # this conversion needed for MLPs

```

```

Y_train = np_utils.to_categorical(y_train, 10)
Y_test = np_utils.to_categorical(y_test, 10)

```

```

print("After converting the output into a vector : ",Y_train[0])

```

```

Class label of first image : 5
After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]

```

```

In [0]: # some model parameters

output_dim = 10
input_dim = X_train.shape[1]

```

```
batch_size = 128
nb_epoch = 20
```

```
In [17]: model_relu = Sequential()
         model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,)))
         model_relu.add(Dense(128, activation='relu'))
         model_relu.add(Dense(output_dim, activation='softmax'))

         model_relu.summary()

         model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

         history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=0)
```

Layer (type)	Output Shape	Param #
dense_4 (Dense)	(None, 512)	401920
dense_5 (Dense)	(None, 128)	65664
dense_6 (Dense)	(None, 10)	1290

```
=====
Total params: 468,874
Trainable params: 468,874
Non-trainable params: 0
```

```
-----
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 7s 121us/step - loss: 0.2420 - acc: 0.9298 - val_loss: 0.1050 - val_acc: 0.9737
Epoch 2/20
60000/60000 [=====] - 7s 114us/step - loss: 0.0859 - acc: 0.9737 - val_loss: 0.0450 - val_acc: 0.9829
Epoch 3/20
60000/60000 [=====] - 7s 114us/step - loss: 0.0558 - acc: 0.9829 - val_loss: 0.0372 - val_acc: 0.9889
Epoch 4/20
60000/60000 [=====] - 7s 116us/step - loss: 0.0372 - acc: 0.9889 - val_loss: 0.0303 - val_acc: 0.9904
Epoch 5/20
60000/60000 [=====] - 7s 117us/step - loss: 0.0303 - acc: 0.9904 - val_loss: 0.0205 - val_acc: 0.9931
Epoch 6/20
60000/60000 [=====] - 7s 113us/step - loss: 0.0205 - acc: 0.9931 - val_loss: 0.0182 - val_acc: 0.9939
Epoch 7/20
60000/60000 [=====] - 7s 113us/step - loss: 0.0182 - acc: 0.9939 - val_loss: 0.0138 - val_acc: 0.9955
Epoch 8/20
60000/60000 [=====] - 7s 113us/step - loss: 0.0138 - acc: 0.9955 - val_loss: 0.0155 - val_acc: 0.9948
Epoch 9/20
60000/60000 [=====] - 7s 113us/step - loss: 0.0155 - acc: 0.9948 - val_loss: 0.0148 - val_acc: 0.9949
Epoch 10/20
60000/60000 [=====] - 7s 114us/step - loss: 0.0148 - acc: 0.9949 - val_loss: 0.0148 - val_acc: 0.9949
```

```

Epoch 11/20
60000/60000 [=====] - 7s 114us/step - loss: 0.0097 - acc: 0.9969 - va
Epoch 12/20
60000/60000 [=====] - 7s 114us/step - loss: 0.0095 - acc: 0.9971 - va
Epoch 13/20
60000/60000 [=====] - 7s 115us/step - loss: 0.0091 - acc: 0.9971 - va
Epoch 14/20
60000/60000 [=====] - 7s 113us/step - loss: 0.0079 - acc: 0.9977 - va
Epoch 15/20
60000/60000 [=====] - 7s 114us/step - loss: 0.0128 - acc: 0.9960 - va
Epoch 16/20
60000/60000 [=====] - 7s 114us/step - loss: 0.0049 - acc: 0.9986 - va
Epoch 17/20
60000/60000 [=====] - 7s 114us/step - loss: 0.0090 - acc: 0.9971 - va
Epoch 18/20
60000/60000 [=====] - 7s 112us/step - loss: 0.0075 - acc: 0.9975 - va
Epoch 19/20
60000/60000 [=====] - 7s 112us/step - loss: 0.0069 - acc: 0.9979 - va
Epoch 20/20
60000/60000 [=====] - 7s 112us/step - loss: 0.0061 - acc: 0.9982 - va

```

```

In [16]: score = model_relu.evaluate(X_test, Y_test, verbose=0)
         print('Test score:', score[0])
         print('Test accuracy:', score[1])

         configure_plotly_browser_state()

         fig,ax = plt.subplots(1,1)
         ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

         # list of epoch numbers
         x = list(range(1,nb_epoch+1))

         # print(history.history.keys())
         # dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
         # history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, validation_data=(X_test, Y_test))

         # we will get val_loss and val_acc only when you pass the paramter validation_data
         # val_loss : validation loss
         # val_acc : validation accuracy

         # loss : training loss
         # acc : train accuracy
         # for each key in history.history we will have a list of length equal to number of epochs

         vy = history.history['val_loss']
         ty = history.history['loss']


```

```
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.09020908048287725

Test accuracy: 0.9836

<IPython.core.display.HTML object>



1 Things keep in mind

<https://stackoverflow.com/questions/47299624/how-to-understand-loss-acc-val-loss-val-acc-in-keras-model-fitting>

Training should be stopped when val_acc stops increasing, otherwise your model will probably overfit. You can use earlystopping callback to stop training.

2 Two Hidden Layers Architecture

```
In [35]: # some model parameters
```

```
output_dim = 10
input_dim = X_train.shape[1]

batch_size = 100
```

```

nb_epoch = 20

# 1st Hidden layer
two_layer_model_relu = Sequential()
two_layer_model_relu.add(Dense(400, activation='relu', input_shape=(input_dim,)))
two_layer_model_relu.add(BatchNormalization())
two_layer_model_relu.add(Dropout(0.5))

# 2nd Hidden layer
two_layer_model_relu.add(Dense(100, activation='relu'))
two_layer_model_relu.add(BatchNormalization())
two_layer_model_relu.add(Dropout(0.5))

two_layer_model_relu.add(Dense(output_dim, activation='softmax'))

two_layer_model_relu.summary()

two_layer_model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = two_layer_model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch)

```

Layer (type)	Output Shape	Param #
dense_32 (Dense)	(None, 400)	314000
batch_normalization_9 (Batch Normalization)	(None, 400)	1600
dropout_3 (Dropout)	(None, 400)	0
dense_33 (Dense)	(None, 100)	40100
batch_normalization_10 (Batch Normalization)	(None, 100)	400
dropout_4 (Dropout)	(None, 100)	0
dense_34 (Dense)	(None, 10)	1010
Total params: 357,110		
Trainable params: 356,110		
Non-trainable params: 1,000		

```

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 11s 180us/step - loss: 0.4294 - acc: 0.8712 - val_loss: 0.4294 - val_acc: 0.8712
Epoch 2/20

```

```

60000/60000 [=====] - 9s 145us/step - loss: 0.2200 - acc: 0.9338 - va
Epoch 3/20
60000/60000 [=====] - 9s 146us/step - loss: 0.1785 - acc: 0.9468 - va
Epoch 4/20
60000/60000 [=====] - 9s 146us/step - loss: 0.1502 - acc: 0.9542 - va
Epoch 5/20
60000/60000 [=====] - 9s 146us/step - loss: 0.1333 - acc: 0.9606 - va
Epoch 6/20
60000/60000 [=====] - 9s 147us/step - loss: 0.1218 - acc: 0.9638 - va
Epoch 7/20
60000/60000 [=====] - 9s 152us/step - loss: 0.1130 - acc: 0.9652 - va
Epoch 8/20
60000/60000 [=====] - 9s 147us/step - loss: 0.1111 - acc: 0.9665 - va
Epoch 9/20
60000/60000 [=====] - 9s 147us/step - loss: 0.1035 - acc: 0.9685 - va
Epoch 10/20
60000/60000 [=====] - 9s 148us/step - loss: 0.0962 - acc: 0.9705 - va
Epoch 11/20
60000/60000 [=====] - 9s 151us/step - loss: 0.0916 - acc: 0.9711 - va
Epoch 12/20
60000/60000 [=====] - 9s 147us/step - loss: 0.0878 - acc: 0.9729 - va
Epoch 13/20
60000/60000 [=====] - 9s 147us/step - loss: 0.0808 - acc: 0.9751 - va
Epoch 14/20
60000/60000 [=====] - 9s 147us/step - loss: 0.0785 - acc: 0.9747 - va
Epoch 15/20
60000/60000 [=====] - 9s 147us/step - loss: 0.0757 - acc: 0.9773 - va
Epoch 16/20
60000/60000 [=====] - 9s 147us/step - loss: 0.0752 - acc: 0.9758 - va
Epoch 17/20
60000/60000 [=====] - 9s 146us/step - loss: 0.0693 - acc: 0.9786 - va
Epoch 18/20
60000/60000 [=====] - 9s 146us/step - loss: 0.0717 - acc: 0.9778 - va
Epoch 19/20
60000/60000 [=====] - 9s 146us/step - loss: 0.0669 - acc: 0.9792 - va
Epoch 20/20
60000/60000 [=====] - 9s 146us/step - loss: 0.0627 - acc: 0.9799 - va

```

```

In [36]: score = two_layer_model_relu.evaluate(X_test, Y_test, verbose=0)
         print('Test score:', score[0])
         print('Test accuracy:', score[1])

         fig,ax = plt.subplots(1,1)
         ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

         # list of epoch numbers
         x = list(range(1,nb_epoch+1))

```



```

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.054605189490824706

Test accuracy: 0.9833



3 Three Hidden Layers Architecture

In [39]: # some model parameters

```

output_dim = 10
input_dim = X_train.shape[1]

batch_size = 100
nb_epoch = 20

# 1st Hidden layer
three_layer_model_relu = Sequential()
three_layer_model_relu.add(Dense(400, activation='relu', input_shape=(input_dim,)))
three_layer_model_relu.add(BatchNormalization())
three_layer_model_relu.add(Dropout(0.5))

# 2nd Hidden layer
three_layer_model_relu.add(Dense(200, activation='relu'))
three_layer_model_relu.add(BatchNormalization())
three_layer_model_relu.add(Dropout(0.5))

# 3rd Hidden layer
three_layer_model_relu.add(Dense(100, activation='relu'))
three_layer_model_relu.add(BatchNormalization())
three_layer_model_relu.add(Dropout(0.5))

three_layer_model_relu.add(Dense(output_dim, activation='softmax'))

three_layer_model_relu.summary()

three_layer_model_relu.compile(optimizer='adam', loss='categorical_crossentropy', met

history = three_layer_model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=1

```

Layer (type)	Output Shape	Param #
dense_40 (Dense)	(None, 400)	314000
batch_normalization_14 (Batch Normalization)	(None, 400)	1600
dropout_8 (Dropout)	(None, 400)	0
dense_41 (Dense)	(None, 200)	80200
batch_normalization_15 (Batch Normalization)	(None, 200)	800
dropout_9 (Dropout)	(None, 200)	0
dense_42 (Dense)	(None, 100)	20100

```

batch_normalization_16 (Batch Normalization) 400
-----
dropout_10 (Dropout) (None, 100) 0
-----
dense_43 (Dense) (None, 10) 1010
=====
Total params: 418,110
Trainable params: 416,710
Non-trainable params: 1,400
-----
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 13s 218us/step - loss: 0.5878 - acc: 0.8226 - val_loss: 0.5878 - val_acc: 0.8226
Epoch 2/20
60000/60000 [=====] - 10s 169us/step - loss: 0.2655 - acc: 0.9224 - val_loss: 0.2655 - val_acc: 0.9224
Epoch 3/20
60000/60000 [=====] - 10s 169us/step - loss: 0.2061 - acc: 0.9400 - val_loss: 0.2061 - val_acc: 0.9400
Epoch 4/20
60000/60000 [=====] - 10s 170us/step - loss: 0.1805 - acc: 0.9487 - val_loss: 0.1805 - val_acc: 0.9487
Epoch 5/20
60000/60000 [=====] - 10s 170us/step - loss: 0.1593 - acc: 0.9543 - val_loss: 0.1593 - val_acc: 0.9543
Epoch 6/20
60000/60000 [=====] - 10s 171us/step - loss: 0.1458 - acc: 0.9577 - val_loss: 0.1458 - val_acc: 0.9577
Epoch 7/20
60000/60000 [=====] - 10s 170us/step - loss: 0.1379 - acc: 0.9598 - val_loss: 0.1379 - val_acc: 0.9598
Epoch 8/20
60000/60000 [=====] - 10s 170us/step - loss: 0.1274 - acc: 0.9626 - val_loss: 0.1274 - val_acc: 0.9626
Epoch 9/20
60000/60000 [=====] - 10s 170us/step - loss: 0.1184 - acc: 0.9653 - val_loss: 0.1184 - val_acc: 0.9653
Epoch 10/20
60000/60000 [=====] - 10s 170us/step - loss: 0.1146 - acc: 0.9657 - val_loss: 0.1146 - val_acc: 0.9657
Epoch 11/20
60000/60000 [=====] - 10s 169us/step - loss: 0.1087 - acc: 0.9679 - val_loss: 0.1087 - val_acc: 0.9679
Epoch 12/20
60000/60000 [=====] - 10s 169us/step - loss: 0.1027 - acc: 0.9693 - val_loss: 0.1027 - val_acc: 0.9693
Epoch 13/20
60000/60000 [=====] - 10s 170us/step - loss: 0.0963 - acc: 0.9721 - val_loss: 0.0963 - val_acc: 0.9721
Epoch 14/20
60000/60000 [=====] - 10s 168us/step - loss: 0.0936 - acc: 0.9727 - val_loss: 0.0936 - val_acc: 0.9727
Epoch 15/20
60000/60000 [=====] - 10s 169us/step - loss: 0.0904 - acc: 0.9733 - val_loss: 0.0904 - val_acc: 0.9733
Epoch 16/20
60000/60000 [=====] - 10s 166us/step - loss: 0.0860 - acc: 0.9741 - val_loss: 0.0860 - val_acc: 0.9741
Epoch 17/20
60000/60000 [=====] - 10s 165us/step - loss: 0.0856 - acc: 0.9750 - val_loss: 0.0856 - val_acc: 0.9750
Epoch 18/20
60000/60000 [=====] - 10s 169us/step - loss: 0.0808 - acc: 0.9762 - val_loss: 0.0808 - val_acc: 0.9762
Epoch 19/20

```

```
60000/60000 [=====] - 10s 170us/step - loss: 0.0779 - acc: 0.9774 - v
Epoch 20/20
60000/60000 [=====] - 10s 169us/step - loss: 0.0743 - acc: 0.9773 - v
```

```
In [40]: score = three_layer_model_relu.evaluate(X_test, Y_test, verbose=0)
        print('Test score:', score[0])
        print('Test accuracy:', score[1])

        fig,ax = plt.subplots(1,1)
        ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

        # list of epoch numbers
        x = list(range(1,nb_epoch+1))

        # print(history.history.keys())
        # dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
        # history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,


        # we will get val_loss and val_acc only when you pass the paramter validation_data
        # val_loss : validation loss
        # val_acc : validation accuracy

        # loss : training loss
        # acc : train accuracy
        # for each key in history.history we will have a list of length equal to number of

        vy = history.history['val_loss']
        ty = history.history['loss']
        plt_dynamic(x, vy, ty, ax)
```

Test score: 0.06033014571936801

Test accuracy: 0.9829



4 Five Hidden Layers Architecture

In [43]: *# some model parameters*

```
output_dim = 10
input_dim = X_train.shape[1]

batch_size = 100
nb_epoch = 20

# 1st Hidden layer
five_layer_model_relu = Sequential()
five_layer_model_relu.add(Dense(500, activation='relu', input_shape=(input_dim,)))
five_layer_model_relu.add(BatchNormalization())
five_layer_model_relu.add(Dropout(0.5))

# 2nd Hidden layer
five_layer_model_relu.add(Dense(350, activation='relu'))
five_layer_model_relu.add(BatchNormalization())
five_layer_model_relu.add(Dropout(0.5))

# 3rd Hidden layer
five_layer_model_relu.add(Dense(200, activation='relu'))
five_layer_model_relu.add(BatchNormalization())
```

```

five_layer_model_relu.add(Dropout(0.5))

# 4th Hidden layer
five_layer_model_relu.add(Dense(100, activation='relu'))
five_layer_model_relu.add(BatchNormalization())
five_layer_model_relu.add(Dropout(0.5))

# 5th Hidden layer
five_layer_model_relu.add(Dense(50, activation='relu'))
five_layer_model_relu.add(BatchNormalization())
five_layer_model_relu.add(Dropout(0.5))

five_layer_model_relu.add(Dense(output_dim, activation='softmax'))

five_layer_model_relu.summary()

five_layer_model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = five_layer_model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=100)

```

Layer (type)	Output Shape	Param #
dense_50 (Dense)	(None, 500)	392500
batch_normalization_22 (Batch Normalization)	(None, 500)	2000
dropout_16 (Dropout)	(None, 500)	0
dense_51 (Dense)	(None, 350)	175350
batch_normalization_23 (Batch Normalization)	(None, 350)	1400
dropout_17 (Dropout)	(None, 350)	0
dense_52 (Dense)	(None, 200)	70200
batch_normalization_24 (Batch Normalization)	(None, 200)	800
dropout_18 (Dropout)	(None, 200)	0
dense_53 (Dense)	(None, 100)	20100
batch_normalization_25 (Batch Normalization)	(None, 100)	400
dropout_19 (Dropout)	(None, 100)	0

dense_54 (Dense)	(None, 50)	5050

batch_normalization_26 (Batch Normalization)	(None, 50)	200

dropout_20 (Dropout)	(None, 50)	0

dense_55 (Dense)	(None, 10)	510
=====		

Total params: 668,510

Trainable params: 666,110

Non-trainable params: 2,400

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 20s 338us/step - loss: 1.0178 - acc: 0.6821 - val_loss: 1.0178 - val_acc: 0.6821

Epoch 2/20

60000/60000 [=====] - 16s 264us/step - loss: 0.3827 - acc: 0.8992 - val_loss: 0.3827 - val_acc: 0.8992

Epoch 3/20

60000/60000 [=====] - 16s 268us/step - loss: 0.2968 - acc: 0.9235 - val_loss: 0.2968 - val_acc: 0.9235

Epoch 4/20

60000/60000 [=====] - 16s 269us/step - loss: 0.2536 - acc: 0.9356 - val_loss: 0.2536 - val_acc: 0.9356

Epoch 5/20

60000/60000 [=====] - 16s 271us/step - loss: 0.2255 - acc: 0.9432 - val_loss: 0.2255 - val_acc: 0.9432

Epoch 6/20

60000/60000 [=====] - 16s 273us/step - loss: 0.1996 - acc: 0.9498 - val_loss: 0.1996 - val_acc: 0.9498

Epoch 7/20

60000/60000 [=====] - 16s 267us/step - loss: 0.1852 - acc: 0.9543 - val_loss: 0.1852 - val_acc: 0.9543

Epoch 8/20

60000/60000 [=====] - 16s 271us/step - loss: 0.1775 - acc: 0.9560 - val_loss: 0.1775 - val_acc: 0.9560

Epoch 9/20

60000/60000 [=====] - 16s 274us/step - loss: 0.1704 - acc: 0.9583 - val_loss: 0.1704 - val_acc: 0.9583

Epoch 10/20

60000/60000 [=====] - 16s 271us/step - loss: 0.1560 - acc: 0.9611 - val_loss: 0.1560 - val_acc: 0.9611

Epoch 11/20

60000/60000 [=====] - 16s 271us/step - loss: 0.1545 - acc: 0.9619 - val_loss: 0.1545 - val_acc: 0.9619

Epoch 12/20

60000/60000 [=====] - 16s 274us/step - loss: 0.1408 - acc: 0.9652 - val_loss: 0.1408 - val_acc: 0.9652

Epoch 13/20

60000/60000 [=====] - 17s 275us/step - loss: 0.1418 - acc: 0.9649 - val_loss: 0.1418 - val_acc: 0.9649

Epoch 14/20

60000/60000 [=====] - 17s 275us/step - loss: 0.1365 - acc: 0.9665 - val_loss: 0.1365 - val_acc: 0.9665

Epoch 15/20

60000/60000 [=====] - 16s 274us/step - loss: 0.1290 - acc: 0.9686 - val_loss: 0.1290 - val_acc: 0.9686

Epoch 16/20

60000/60000 [=====] - 16s 272us/step - loss: 0.1274 - acc: 0.9675 - val_loss: 0.1274 - val_acc: 0.9675

Epoch 17/20

60000/60000 [=====] - 16s 269us/step - loss: 0.1186 - acc: 0.9705 - val_loss: 0.1186 - val_acc: 0.9705

Epoch 18/20

```

60000/60000 [=====] - 16s 267us/step - loss: 0.1209 - acc: 0.9703 - v
Epoch 19/20
60000/60000 [=====] - 16s 269us/step - loss: 0.1130 - acc: 0.9727 - v
Epoch 20/20
60000/60000 [=====] - 16s 271us/step - loss: 0.1089 - acc: 0.9729 - v

```

```

In [44]: score = five_layer_model_relu.evaluate(X_test, Y_test, verbose=0)
         print('Test score:', score[0])
         print('Test accuracy:', score[1])

         fig,ax = plt.subplots(1,1)
         ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

         # list of epoch numbers
         x = list(range(1,nb_epoch+1))

         # print(history.history.keys())
         # dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
         # history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, validation_data=(X_test, Y_test))

         # we will get val_loss and val_acc only when you pass the paramter validation_data
         # val_loss : validation loss
         # val_acc : validation accuracy


         # loss : training loss
         # acc : train accuracy
         # for each key in history.history we will have a list of length equal to number of epochs

         vy = history.history['val_loss']
         ty = history.history['loss']
         plt_dynamic(x, vy, ty, ax)

```

Test score: 0.07435557465390302

Test accuracy: 0.9829



5 Conclusion

```
In [21]: import pandas as pd
         from prettytable import PrettyTable
```

```
bold = '\033[1m'
end = '\033[0m'
```

```
print(bold+'\t\t\t\t Keras ' +end)
print('\n')
```

```
print('In Two hidden layers we have used, neuron architecture is '+bold+str(400)+'\n')
print('In Three hidden layers we have used,neuron architecture is '+bold+str(400)+'\n')
print('In Five hidden layers we have used, neuron architecture is '+bold+str(500)+'\n')
```

```
x = PrettyTable()
x.field_names = ['Metric','Two Hidden Layer','Three Hidden Layer', 'Five Hidden Layer']

x.add_row(["Train Accuracy ", 0.9799,0.9773,0.9729])
x.add_row(["Train Loss ", 0.0627,0.0743,0.1089])
```

```

x.add_row(["Validation Accuracy ",0.9833,0.9829,0.9829])
x.add_row(["Validation Loss ", 0.0546,0.0603,0.0744])

x.add_row(["Test Accuracy ",0.9833,0.9829,0.9829])
x.add_row(["Test Loss ", 0.054605,0.06033,0.074355])

print('\n')
print(x)

```

Keras

In Two hidden layers we have used, neuron architecture is 400,100
 In Three hidden layers we have used,neuron architecture is 400,200,100
 In Five hidden layers we have used, neuron architecture is 500,350,200,100,50

Metric	Two Hidden Layer	Three Hidden Layer	Five Hidden Layer
Train Accuracy	0.9799	0.9773	0.9729
Train Loss	0.0627	0.0743	0.1089
Validation Accuracy	0.9833	0.9829	0.9829
Validation Loss	0.0546	0.0603	0.0744
Test Accuracy	0.9833	0.9829	0.9829
Test Loss	0.054605	0.06033	0.074355

In []: